

Tuesday, May 13, Practical Session II: Neo4j

Ryan Horne 

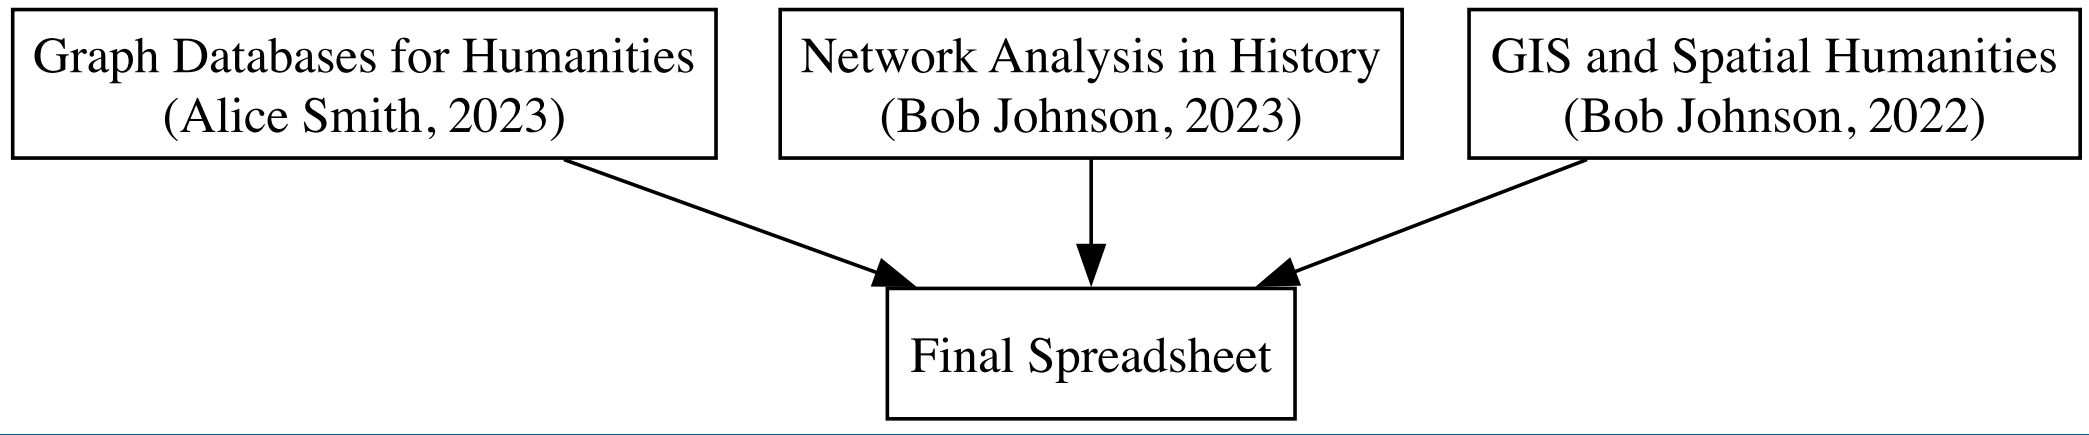
Advanced Research Computing, UCLA
ryan.matthew.horne@gmail.com

Agenda

- From data to databases
- Fundamentals of graph databases
- Neo4j operations

Tabular Data

- Familiar to most of us
- Excell, Google Sheets, etc
- *Rows and Columns*



ID	Title	Author	Year
1	Graph Databases for Humanities	Alice Smith	2023
2	Network Analysis in History	Bob Johnson	2023
3	GIS and Spatial Humanities	Bob Johnson	2022

Some Issues

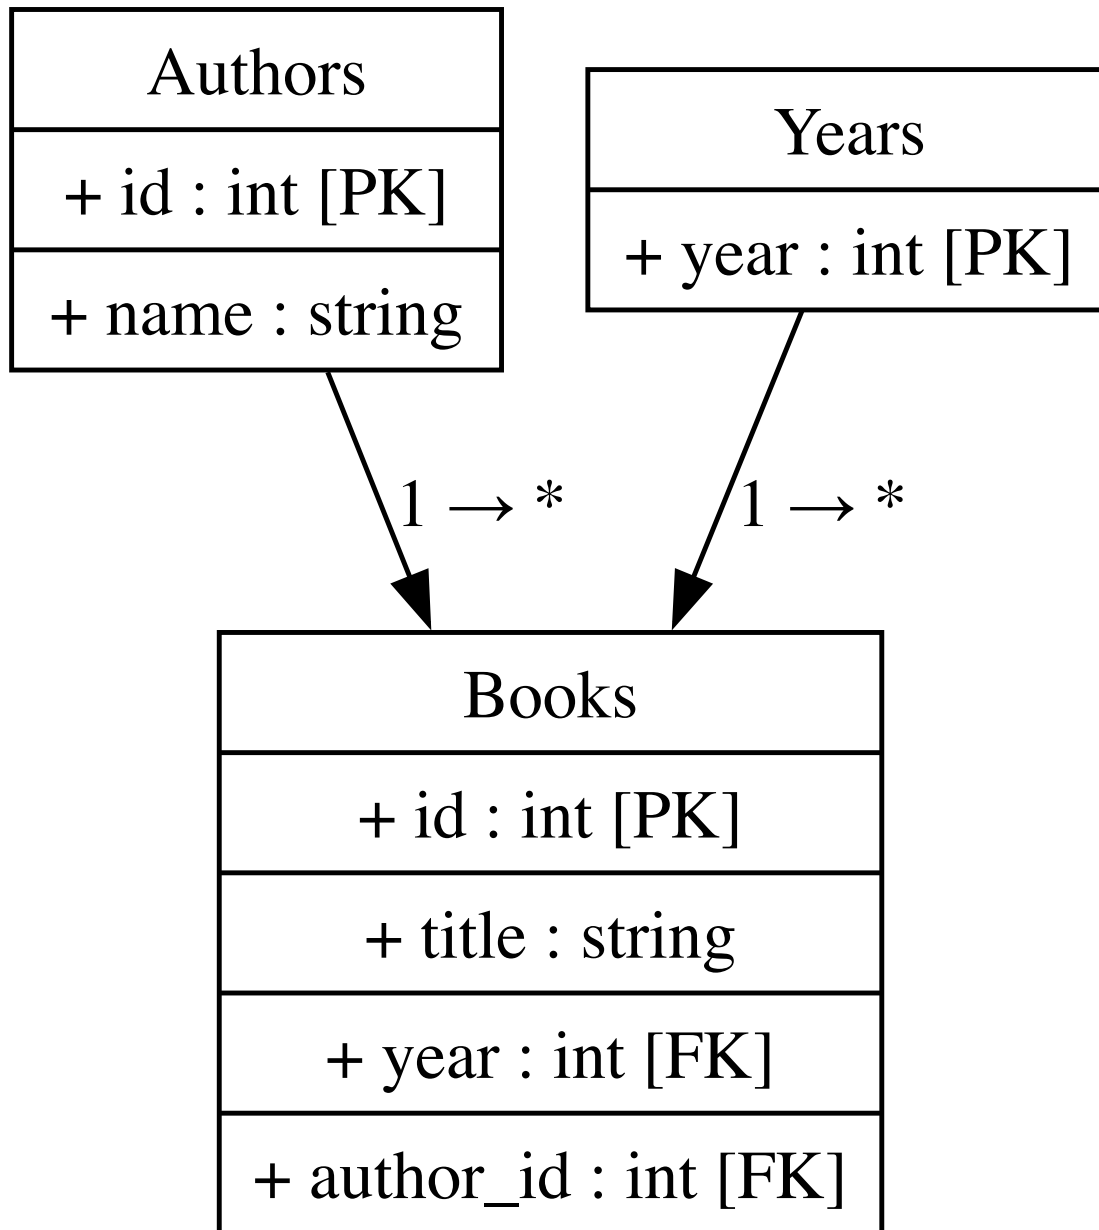
- Not a very flexible data model
- Data redundancy
- Cascading changes
- Not an easy way to link different tables and files

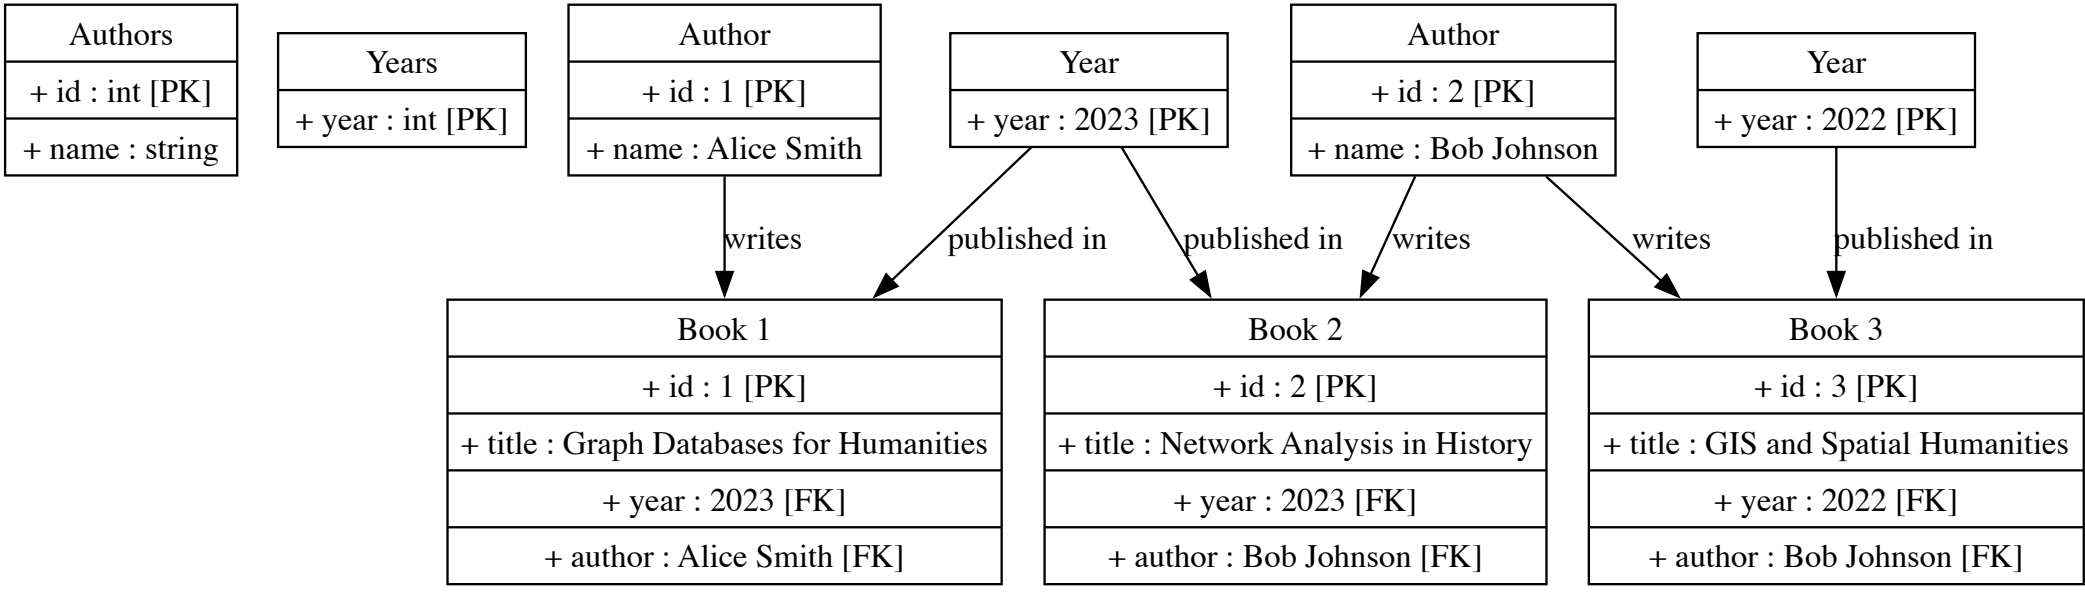
Relational Model

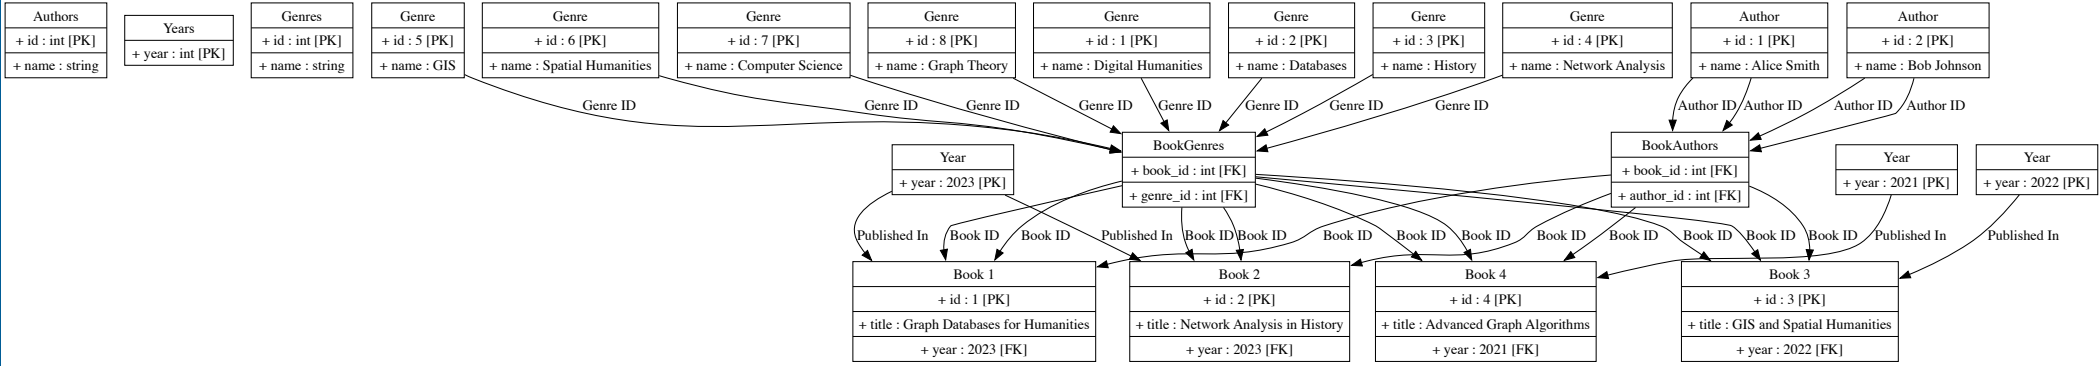
- A full decoupling of the logic of organizing data from the physical structures that house it
- Proposed in 70's; embraced by the 80's
- Schema driven
 - Structure predefined
- Standardized language through SQL
- Same idea as spreadsheets, but we can now link tables through *keys*

Keys to the Kingdom

- *Primary key*: Unique identifier (like ID) which is not repeated in that column
- *Foreign key*: A column whose value is the *primary key* of another row (entity)
- You can have multiple foreign keys in the same table; a linking table can have two columns with the values occurring multiple times
 - It would still have its *own* primary key though!







Problems with the Relational Model

- Structure known *before* data is added
- Complex queries to represent many to 1 and many to many situations
- Despite the name....not really about relations!

A Sequel to SQL: NoSQL

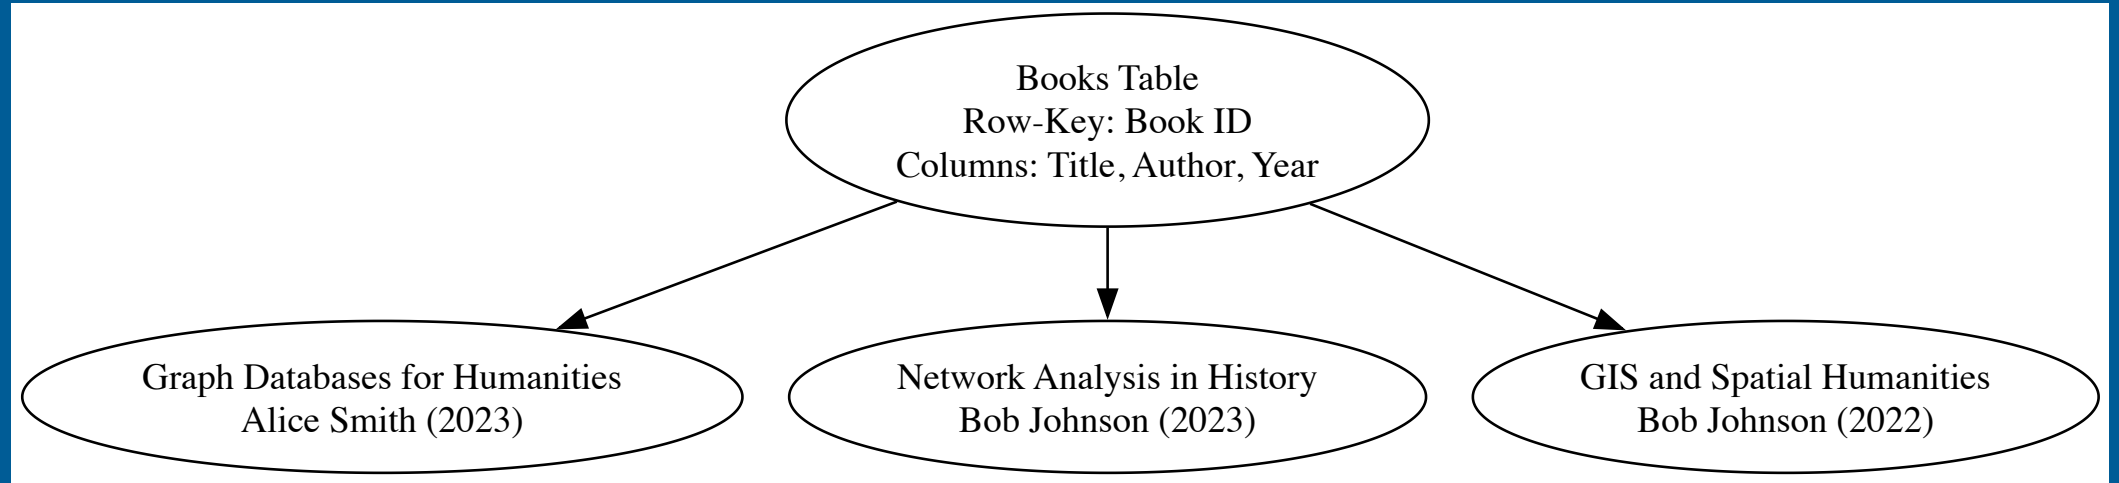
- NoSQL = Not Only SQL
- Umbrella term for databases that:
 - Do not (only) use SQL
 - May not have a fixed schema
 - Might Not use tabular structure
- Key Part: *A flexible schema*

Main Categories

- Wide Column Stores / Column-Family
- Key-Value
- Document
- Graph

Wide Column Stores / Column-Family

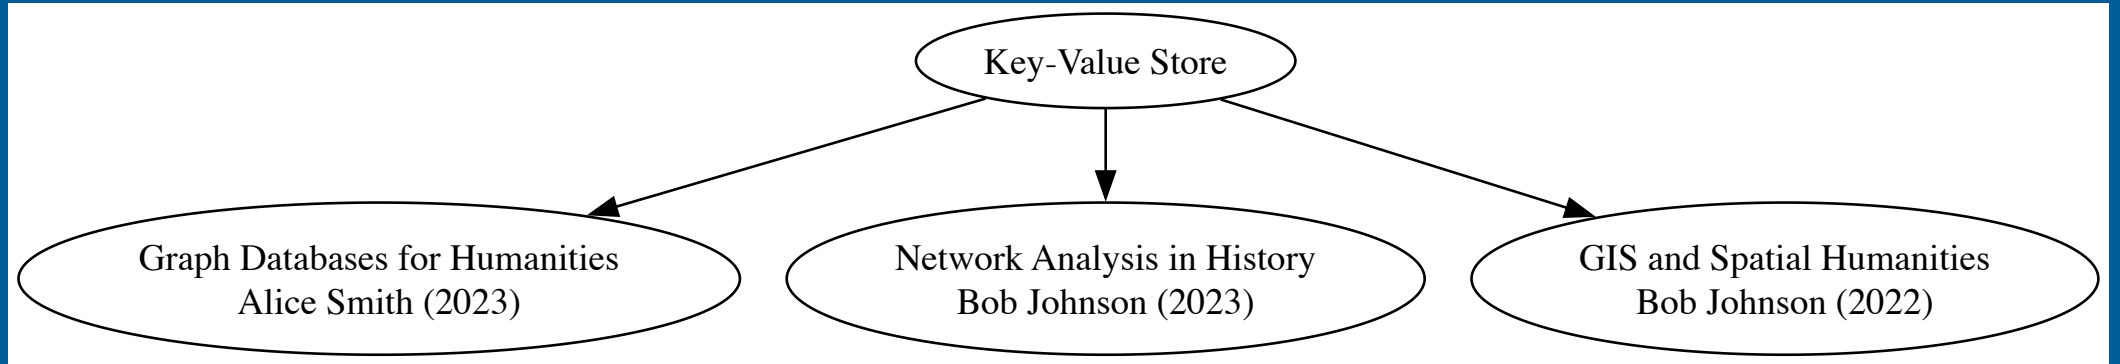
- Rows and Columns in a tabular format
- Row key identifies each record
- Columns grouped into families
- Number of columns for each row is not fixed



Feature	Column-Family Database
Structure	Columns grouped into families
Schema	Flexible or semi-structured schema
Query Language	Query language (e.g., CQL in Cassandra)
Relationships	Limited or no relationships
Use Cases	Time-series data, IoT, logs, and analytics

Key-Value

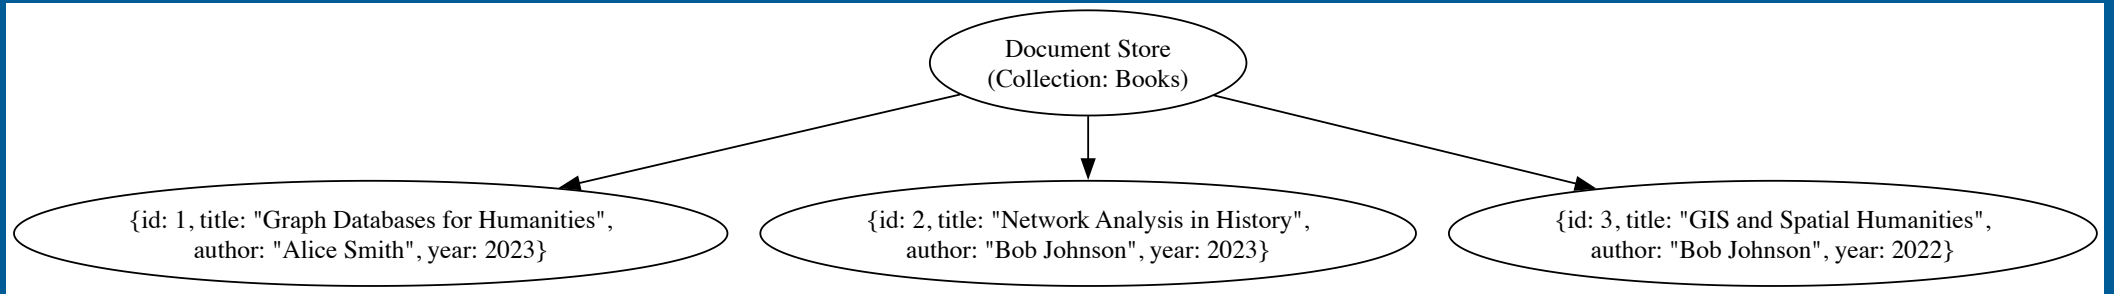
- Key-value pairs, data structures where the *key* is used to identify an entity that is associated with a *value*
- Very Flexible
- Terrible at storing relationships



Feature	Key-Value Database
Structure	Key-value pairs
Schema	Schema-free
Query Language	Simple API (get, put, delete)
Relationships	No relationships by default
Use Cases	Caching, session storage, configuration

Document Model

- Extends the idea of Key:Value pairs by introducing the paradigm of storing data as *documents*
- Typically in a JSON, BSON, or XML format
- Each document can have its own structure and schema as required
- Again not great at relationships



Feature	Document Database
Structure	JSON or BSON documents
Schema	Schema-free or dynamic schema
Query Language	JSON-based query API (e.g., MongoDB Query Language)
Relationships	Embedded documents or references
Scalability	Horizontal scaling
Use Cases	Content management, catalogs, user profiles

On to Graph Databases!

What is a Graph Database?

- Key components: nodes, relationships, and properties
- Built on a network data model

Basic Idea Revisited

- Networks are a collection of entities
- At least some are *linked*
- All kinds of subject domains
- Very flexible definition

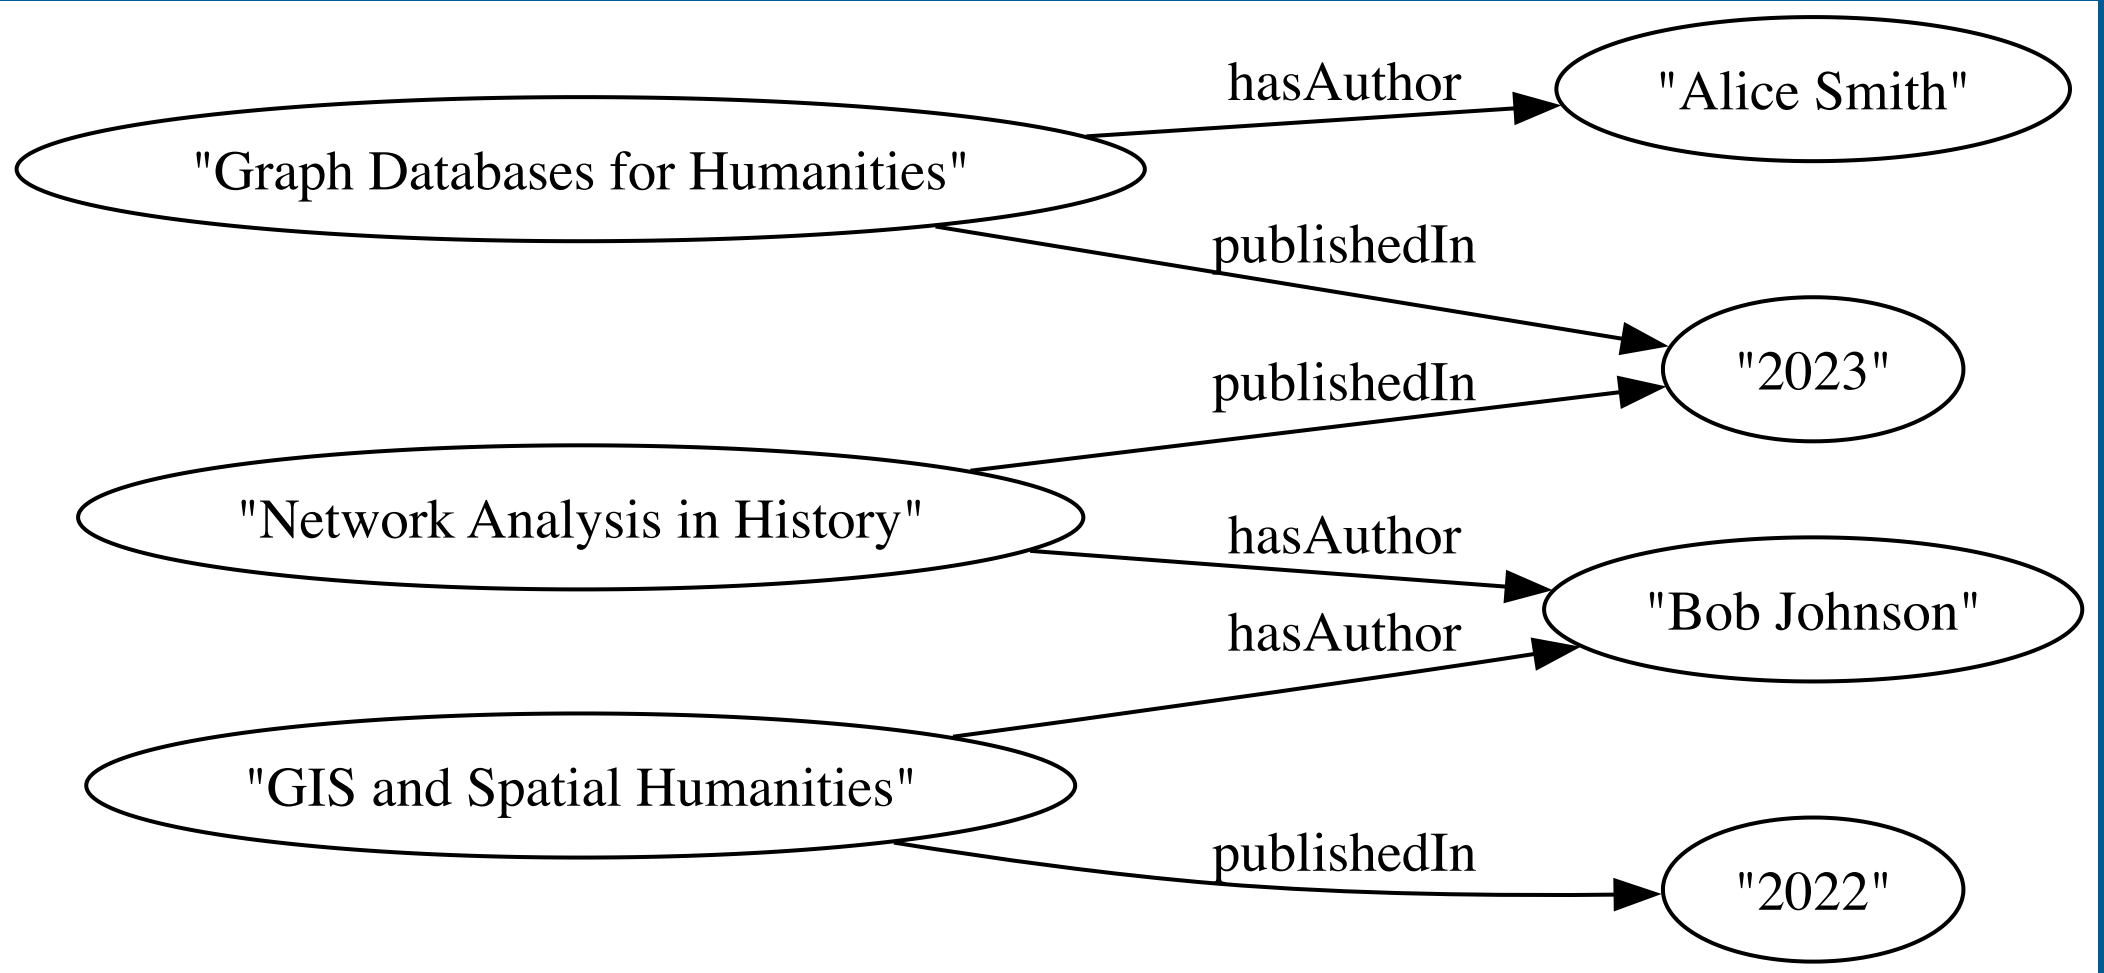
Types of Graph Databases

There are two main types:

- RDF / Triple Store
- Linked Property
- Also Wikidata/Wikibase which is not a *true* graph database

RDF/Triple Store Graphs

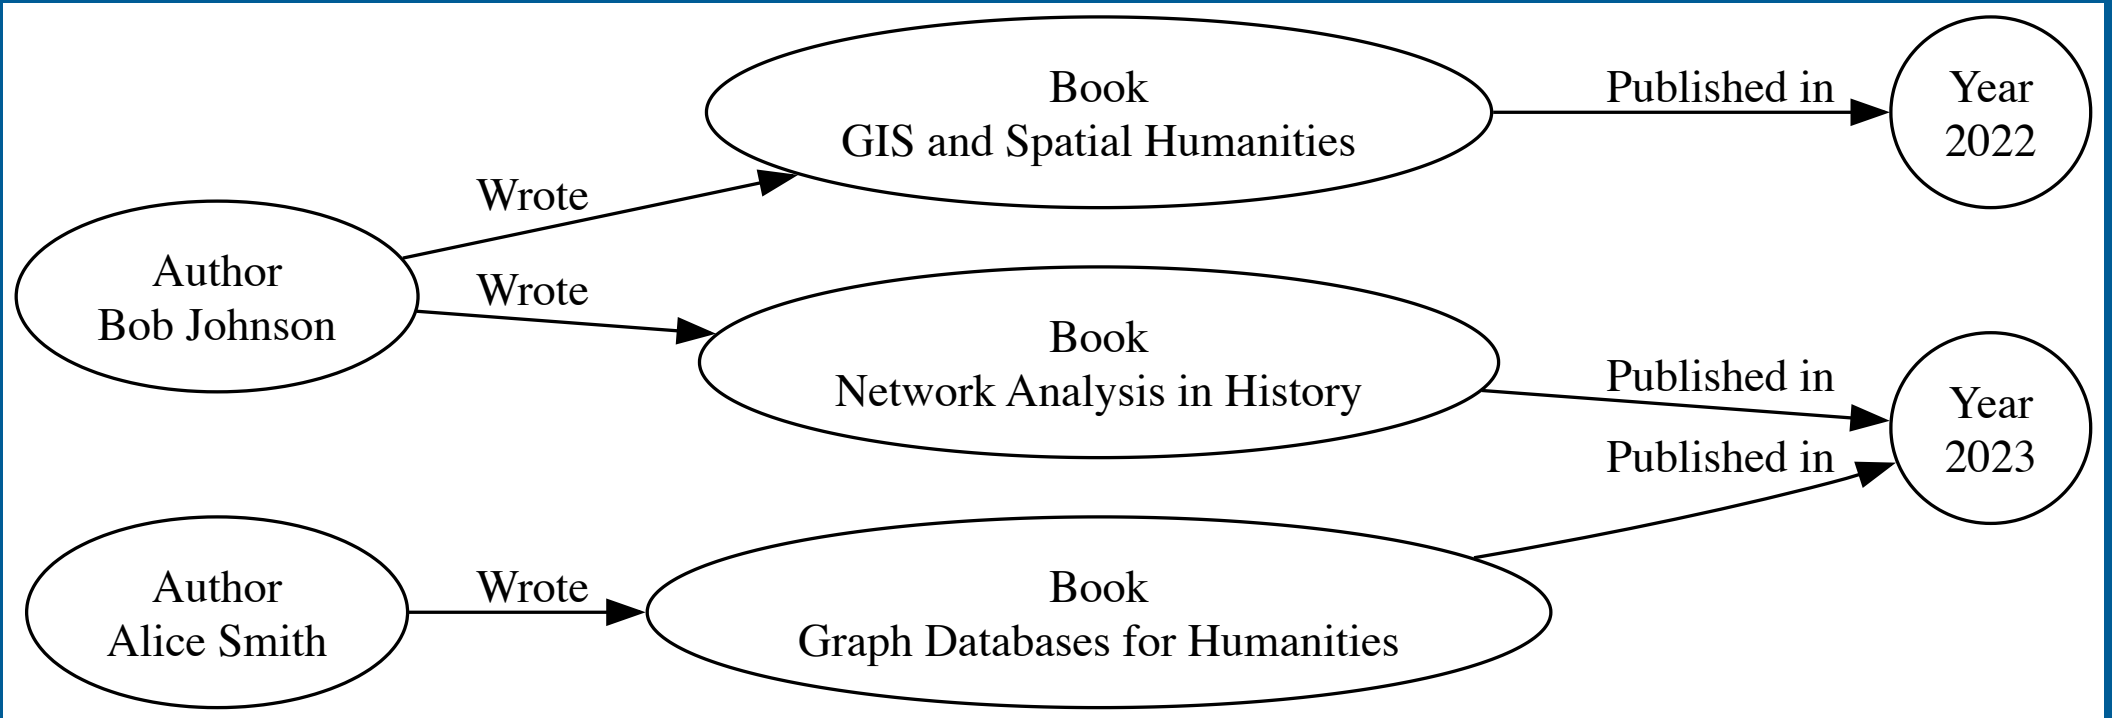
- Backbone of the semantic web and almost all Linked Open Data
- Key technology in Big Data
- RDF is a data exchange standard from W3C
 - Data model of subject (the entity being described)- predicate-(the relationship between the subject and the object) and the object (the target or value of the statement).
 - `<This person><wrote><that book>`



Feature	RDF Graph Database
Structure	Triples (subject-predicate-object)
Schema	Schema-optional (RDFS/OWL for semantics)
Query Language	SPARQL (Semantic query language)
Relationships	Relationships via triples (semantic links)
Use Cases	Linked data, semantic web, ontologies

Labeled Property Graphs

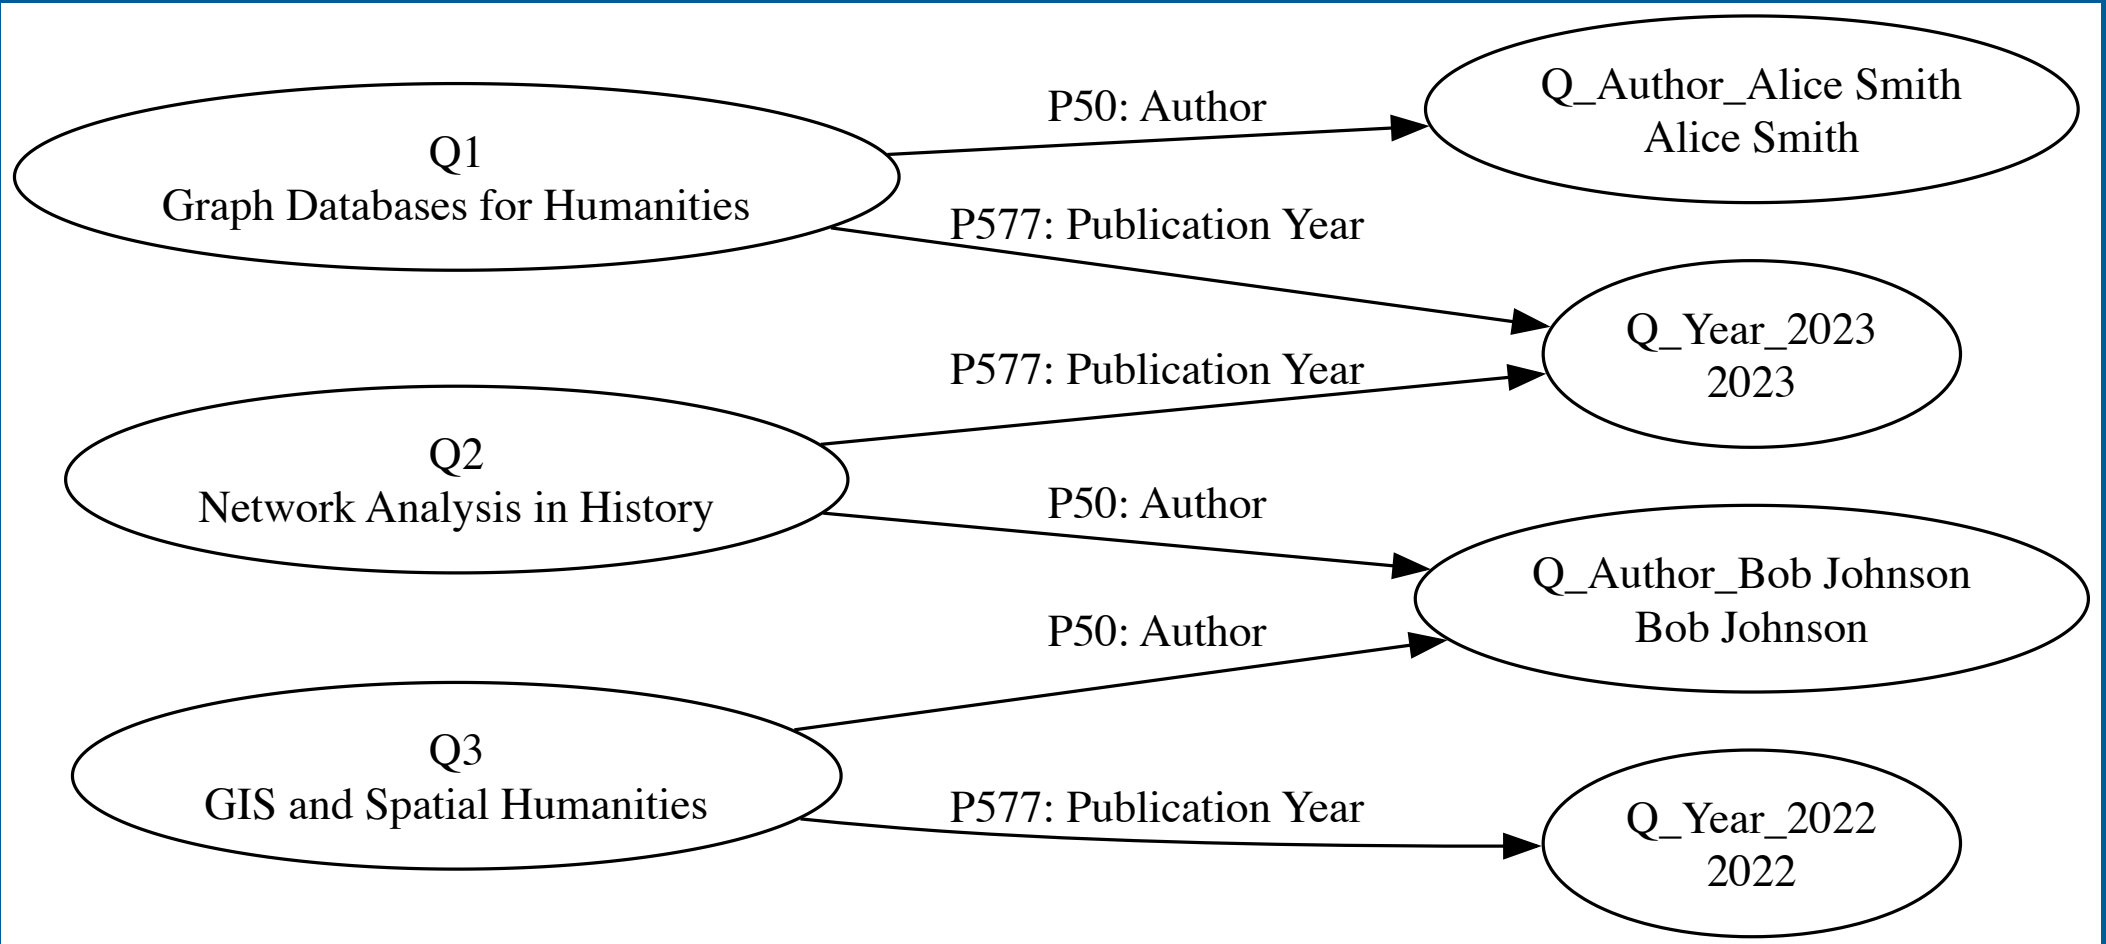
- Allow labeling for nodes (sometimes called vertices or actors) and edges, along with any number of properties for each
- Relationships are inherently directed, with a start and an end node
- Nodes and edges are primitive types; no need to linking tables
- No industry standard querying language
- Neo4j is currently the most popular implementation



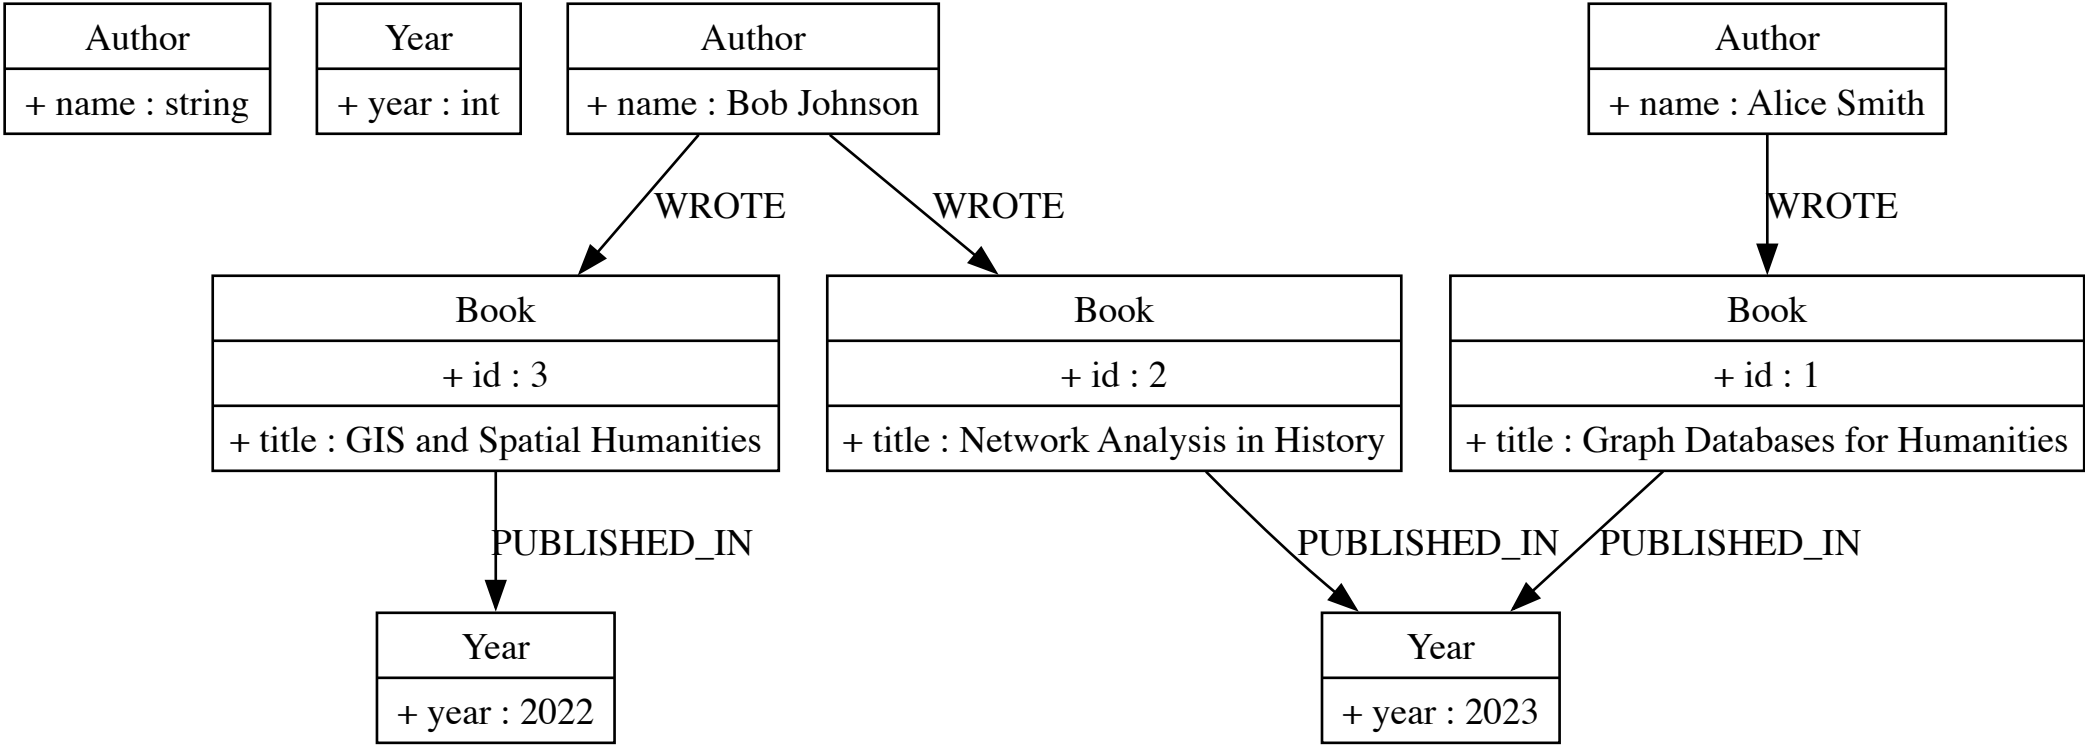
Feature	Property Graph Database
Structure	Nodes, edges, and properties
Schema	Schema-optional or dynamic
Query Language	Graph-specific (e.g., Cypher, Gremlin)
Relationships	Native relationships with edges
Use Cases	Social networks, recommendation engines, fraud detection

A Hybrid Approach: The Case of Wikidata/ Wikibase

- Not a “true” graph database
 - Underlying database is relational, but uses graph-like structures
- Entity-Relationship model
- *Items* have *properties* which can connect to other items
- Primarily intended to work with semantic web so uses RDF like structures and encoding

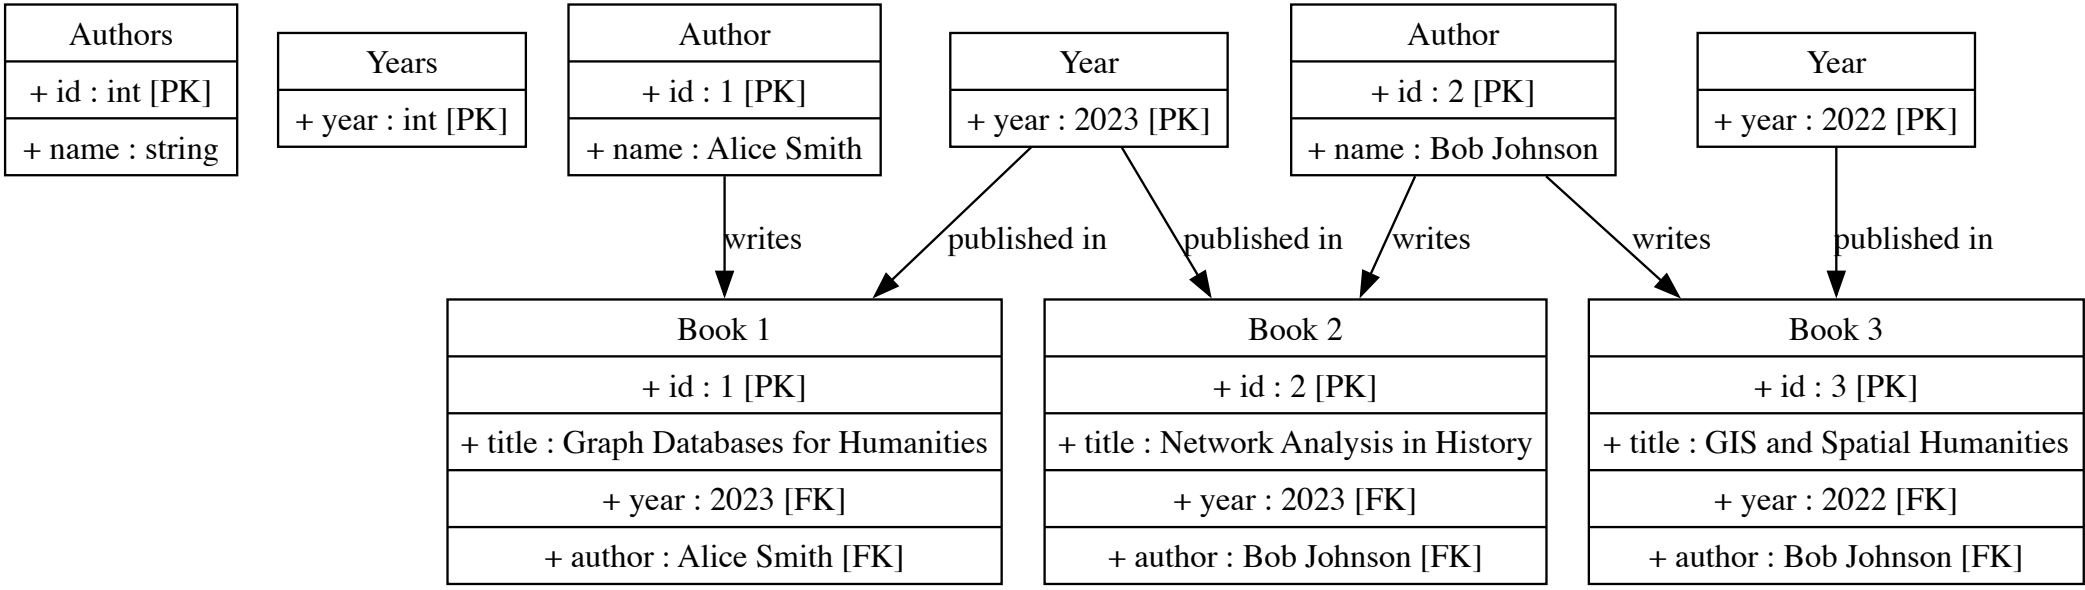



Feature	Wikibase	Graph Database	SQL Database
Structure	Entities (items, properties, values), statements	Nodes, edges, and properties	Tables with rows and columns
Schema	Schema-optional (but predefined entities like items, properties)	Schema-optional or dynamic	Predefined schema required
Query Language	SPARQL, Wikibase Query Service (WQS)	Graph-specific (e.g., Cypher, Gremlin)	SQL
Relationships	Relationships via statements (subject, predicate, object)	Native relationships with edges	Relationships via foreign keys

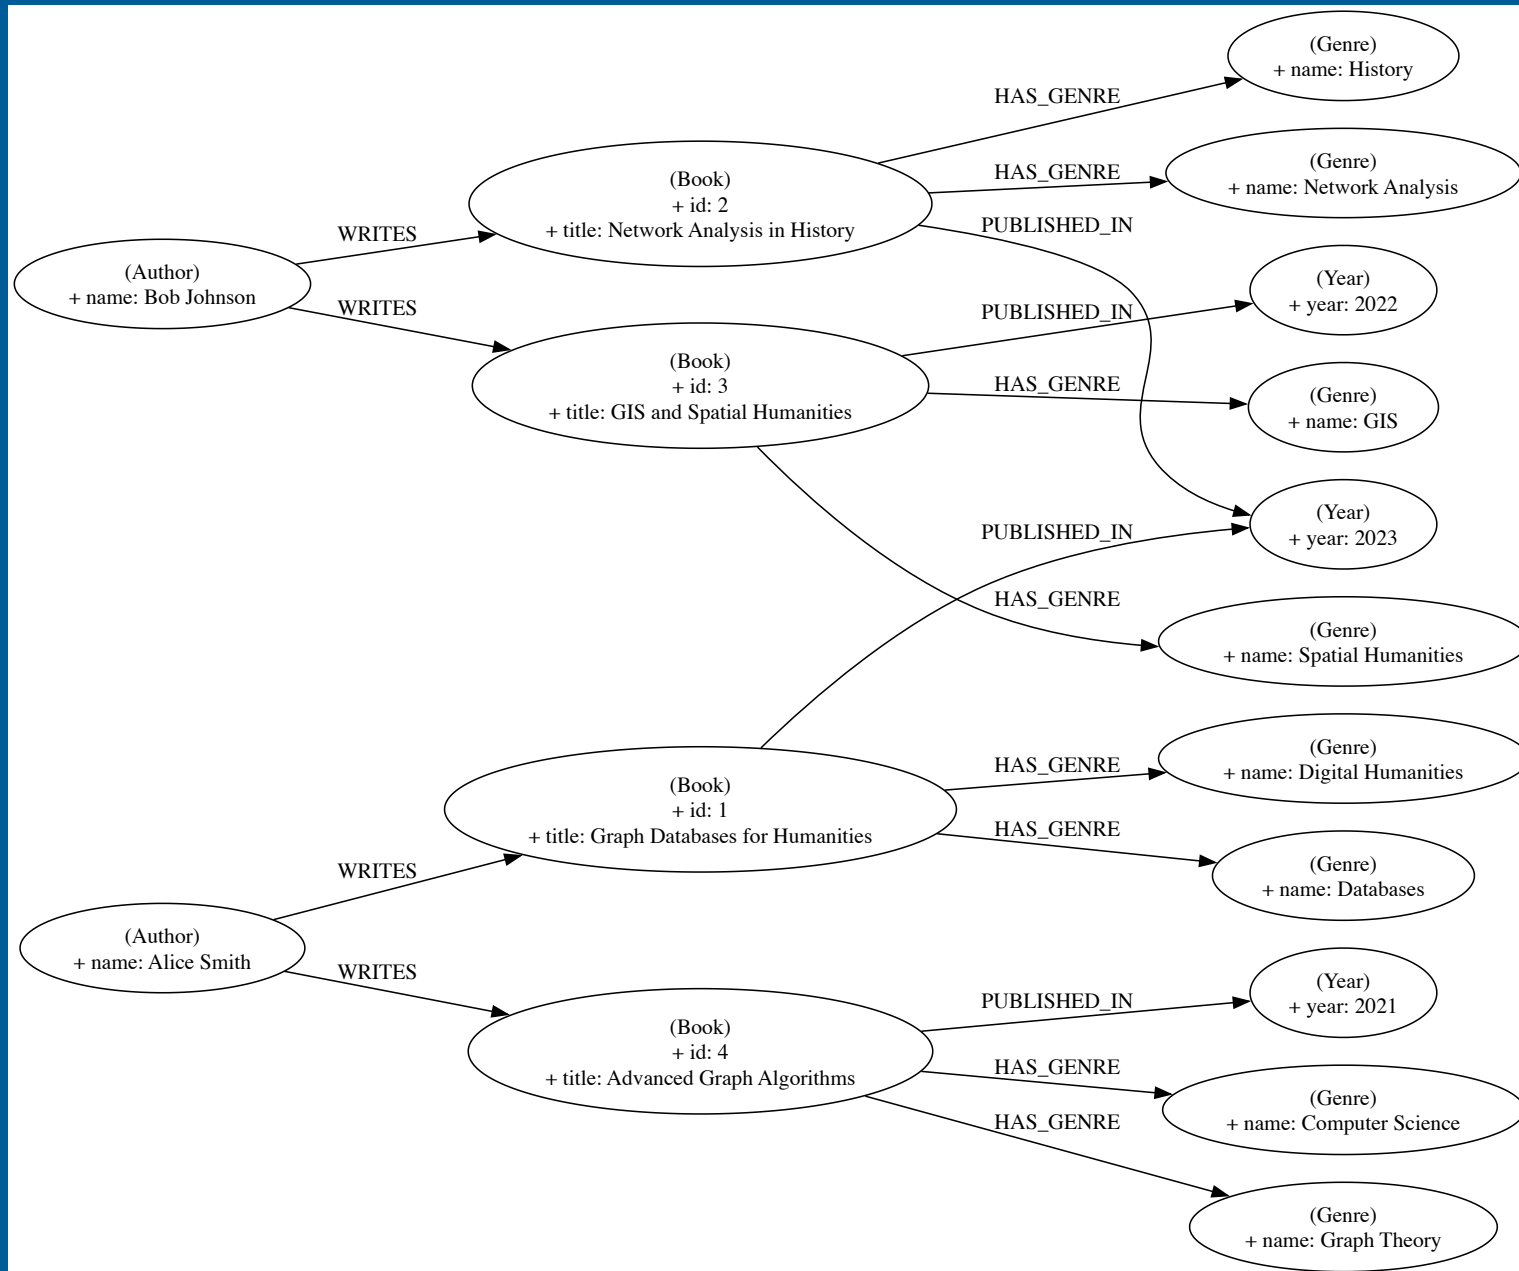


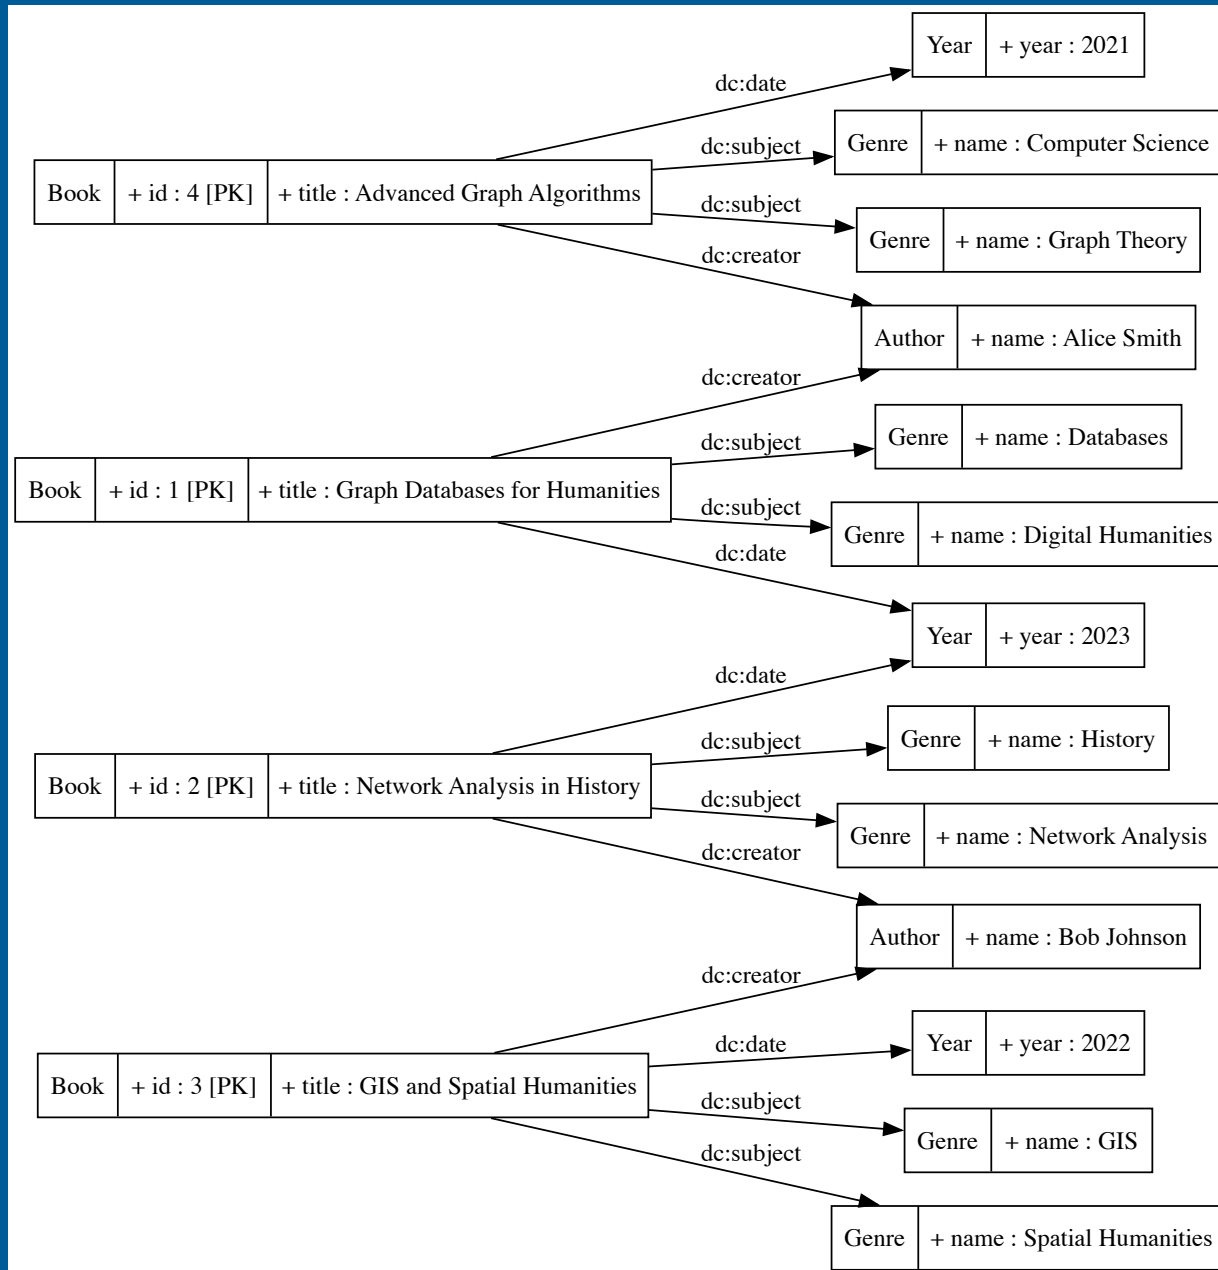
When to use a Graph Database?

- Modeling complex relationships (e.g., networks of people, places, and texts).
- Supporting research questions related to connectivity, lineage, and influence.









Why a Graph Database?

- Operations directly on a graph schema
 - Constraints can be used
- Relationships are seamless and can be changed when necessary
- Multi-layer intersecting domains and fields
- Vocabularies and ontologies to help with data modeling
- Focus on *connections*
- This is a *network* graph!!

Why a Graph Database?

- Models the thought process
- Each node can contain just what it needs and no more
- Ability to adapt and change to the data schema on the fly
- Already designed to facilitate networks and network connections

Why a Graph Database?

- Outputs / inputs can be in different models
 - Tables are very popular
 - Can write scripts to transform lists / json / other structures
 - Programs like Gephi can be used as inputs with the right software
- Extensive support in python, quarto, etc for manipulation and presentation

Issues

- No standard language (although it is being worked on)
- Still not well-known in some fields of study
- Need to think of your data as a network
- Could be a lack of experience with digital / computational studies

Neo4j and Cypher

- Cypher is a graph query language that is used to query the Neo4j Database.
- It is proprietary, but there are open source movements
- Also heavily influencing the development of GQL, which will hopefully standardize things

Neo4j Deployment Options

- Neo4j Desktop: Ideal for local research and development
- Neo4j Aura (Cloud): Free and paid hosted options
- Docker: Advanced users and integration
- APOC & Graph Data Science: Libraries for extended functionality

Hands-On Neo4j SandBox

We are using the movies database

- Click on the database icon
- Look at the sections; we have:
 - node labels
 - relationship types
 - property keys
- Click on the **Person** button

Hands-On Neo4j SandBox

- In writing a Cypher query, a node is enclosed between a parenthesis
 - `(p:Person)`
 - `p` is a variable (what we are calling it) and *Person* is the node
 - Also needs a **RETURN** statement
 - We are using **MATCH** that specifies the patterns we will use in the data
 - Often with **WHERE** clause

Write something!

Can you think of a way to write the following query:

- Load movies into variable **m**
- Where the released property > 1999
- Now just return the *count*

Relationships

- Enclosed in square brackets
 - `[w:WORKS_FOR]`
 - w is a variable
 - WORKS_FOR is the type of relationship

What do you think this does?

```
1 MATCH (p:Person)-[d:DIRECTED]-(m:Movie)
2 WHERE m.released > 2010
3 RETURN p,d,m
```

Your Turn

Query to get all the people who acted in a movie that was released after 2010.

Another question

What is the difference between these two queries?

```
1 MATCH (p:Person)
2 RETURN p
3 LIMIT 20
```

```
1 MATCH (n)
2 RETURN n
3 LIMIT 20
```

MATCH vs MERGE

- MATCH assumes that a node with the specified query exists
- MERGE adds the node if it does not

Match

```
1 MATCH (p:Person)
2 WHERE p.name = "John Doe"
3 SET p.personstatus = 'found'
4 RETURN p
```


Merge

```
1  
2 MERGE (p:Person {name: "John Doe"})  
3 ON MATCH SET p.personstatus = 'found'  
4 RETURN p
```

Create a Relationship Between Nodes

```
1 MATCH (p:Person), (m:Movie)
2 WHERE p.name = "Tom Hanks" AND m.title = "Cloud Atlas"
3
4 CREATE (p)-[w:WATCHED]->(m)
5 RETURN type(w)
```

Wait, what is this?

- Did you see the “->” above?
- What is going on?

Directed vs. Undirected Graph

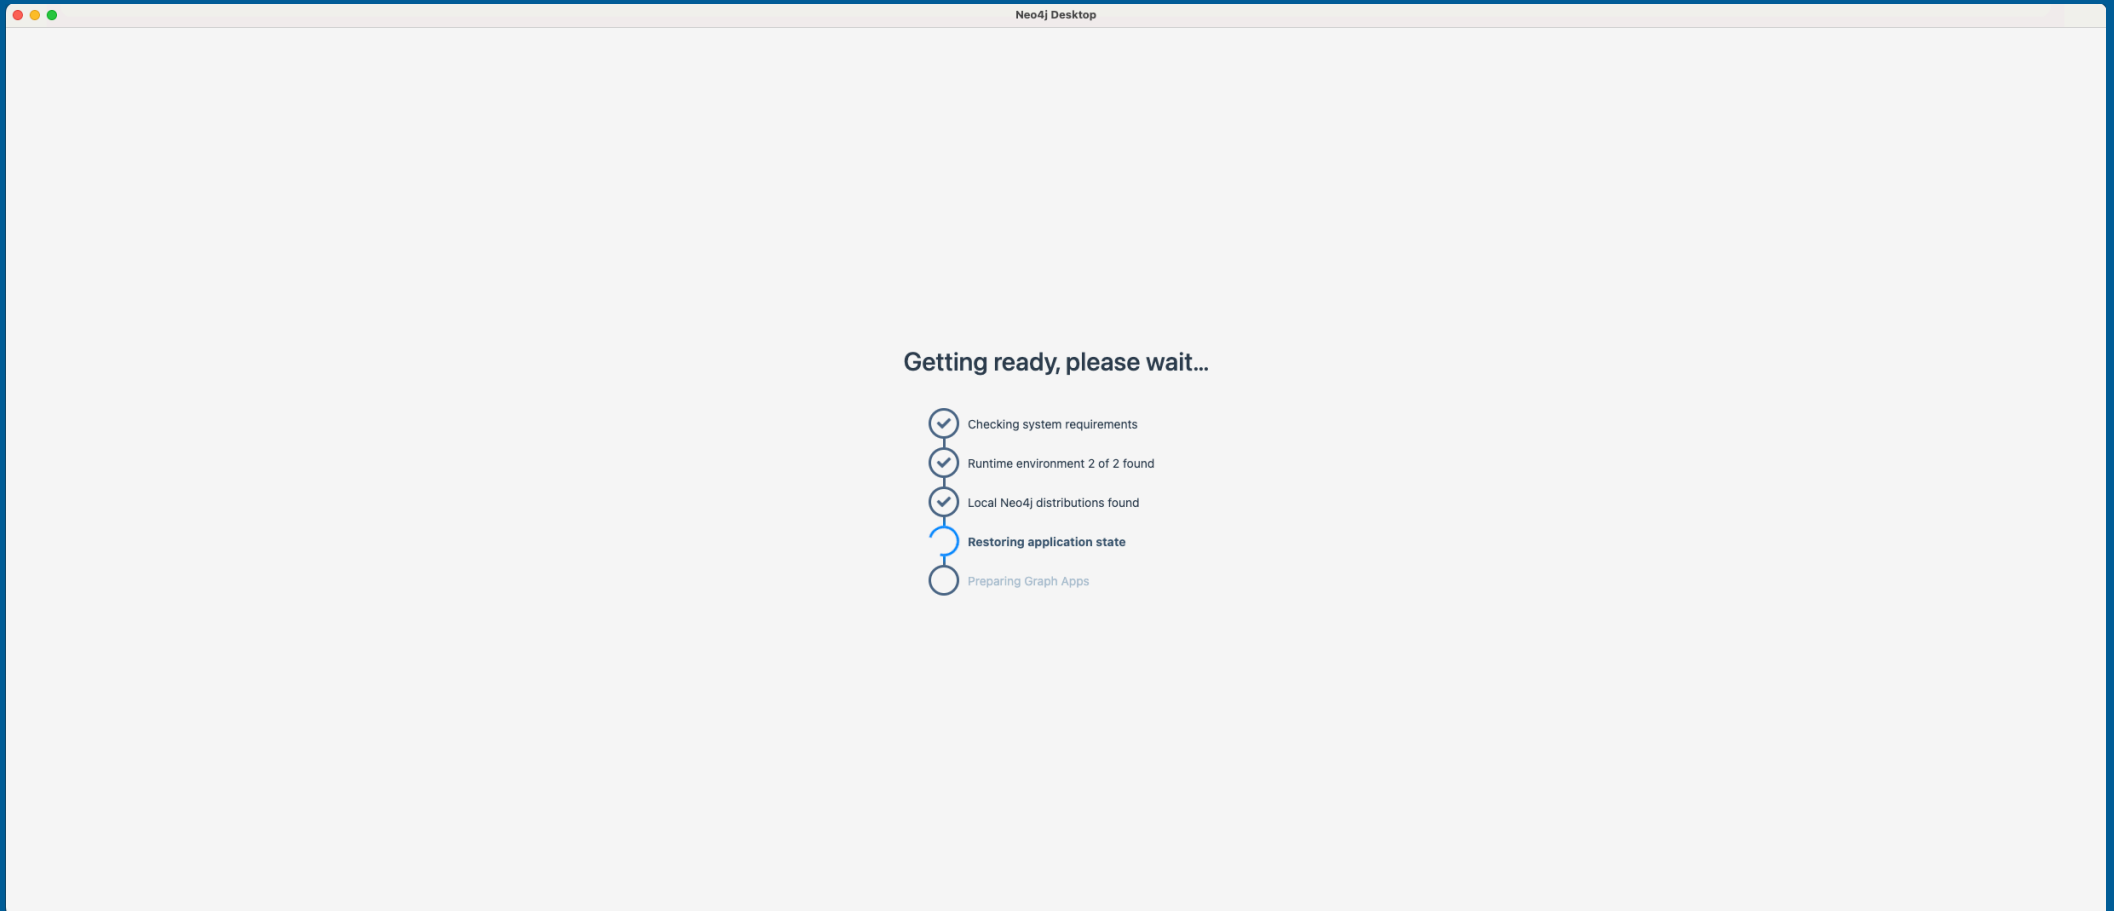
- We can specify the *direction* of the relationship
 - or just leave out the arrow for undirected connections
- Why might this be important?
- When do we not care (or that direction does not make sense?)

Find all People With any Relation to a Node

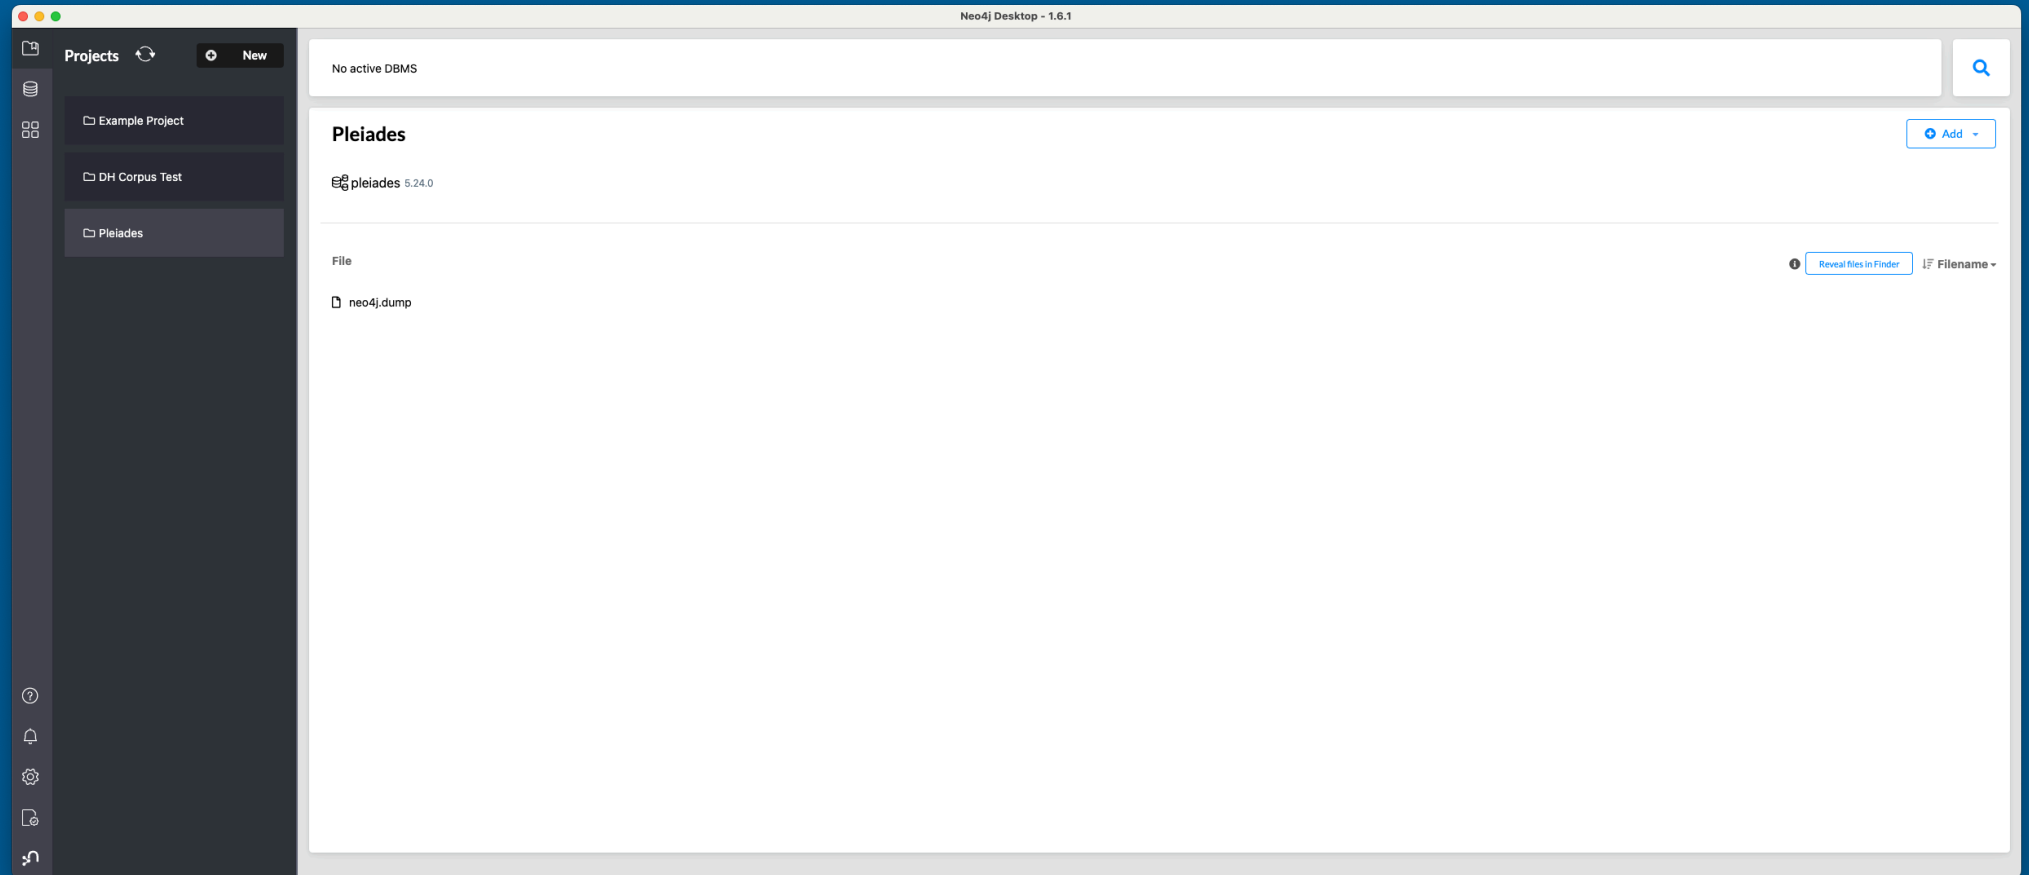
```
1  
2 MATCH (p:Person)-[relatedTo]-(m:Movie {title: "Cloud Atlas"})  
3 RETURN p.name, type(relatedTo)  
4
```

Desktop Neo4j

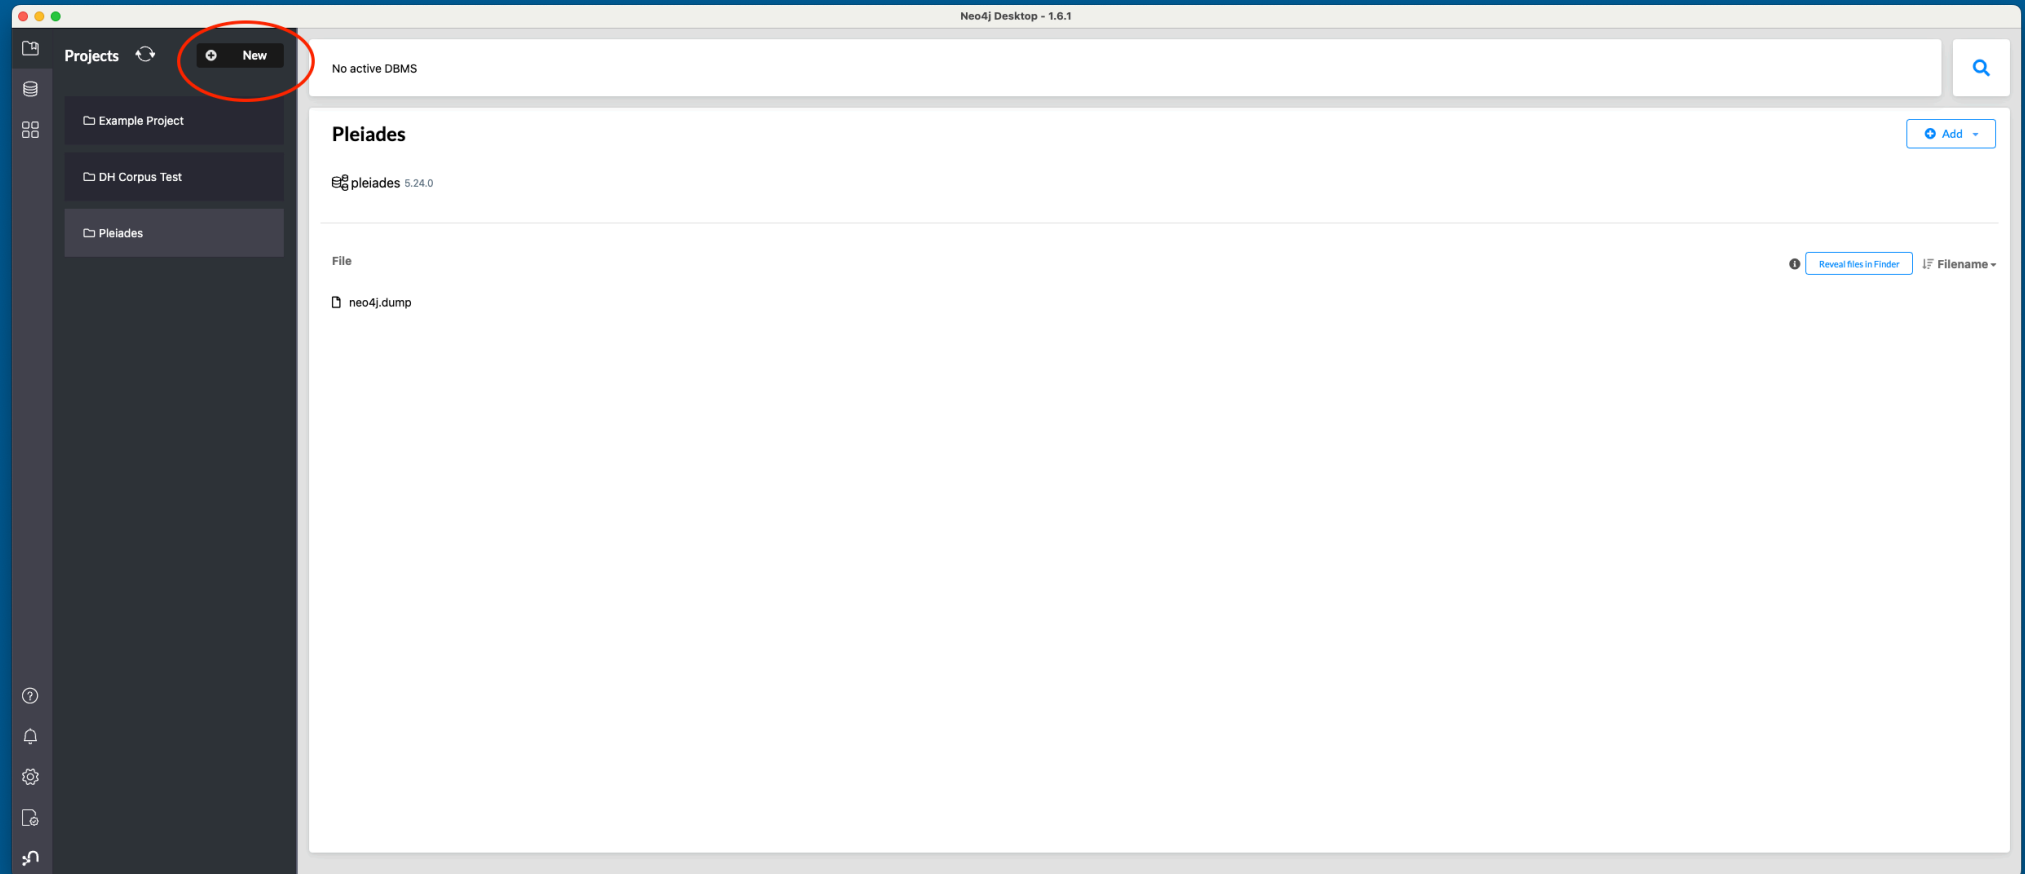
After we install the application, we should see something like this:



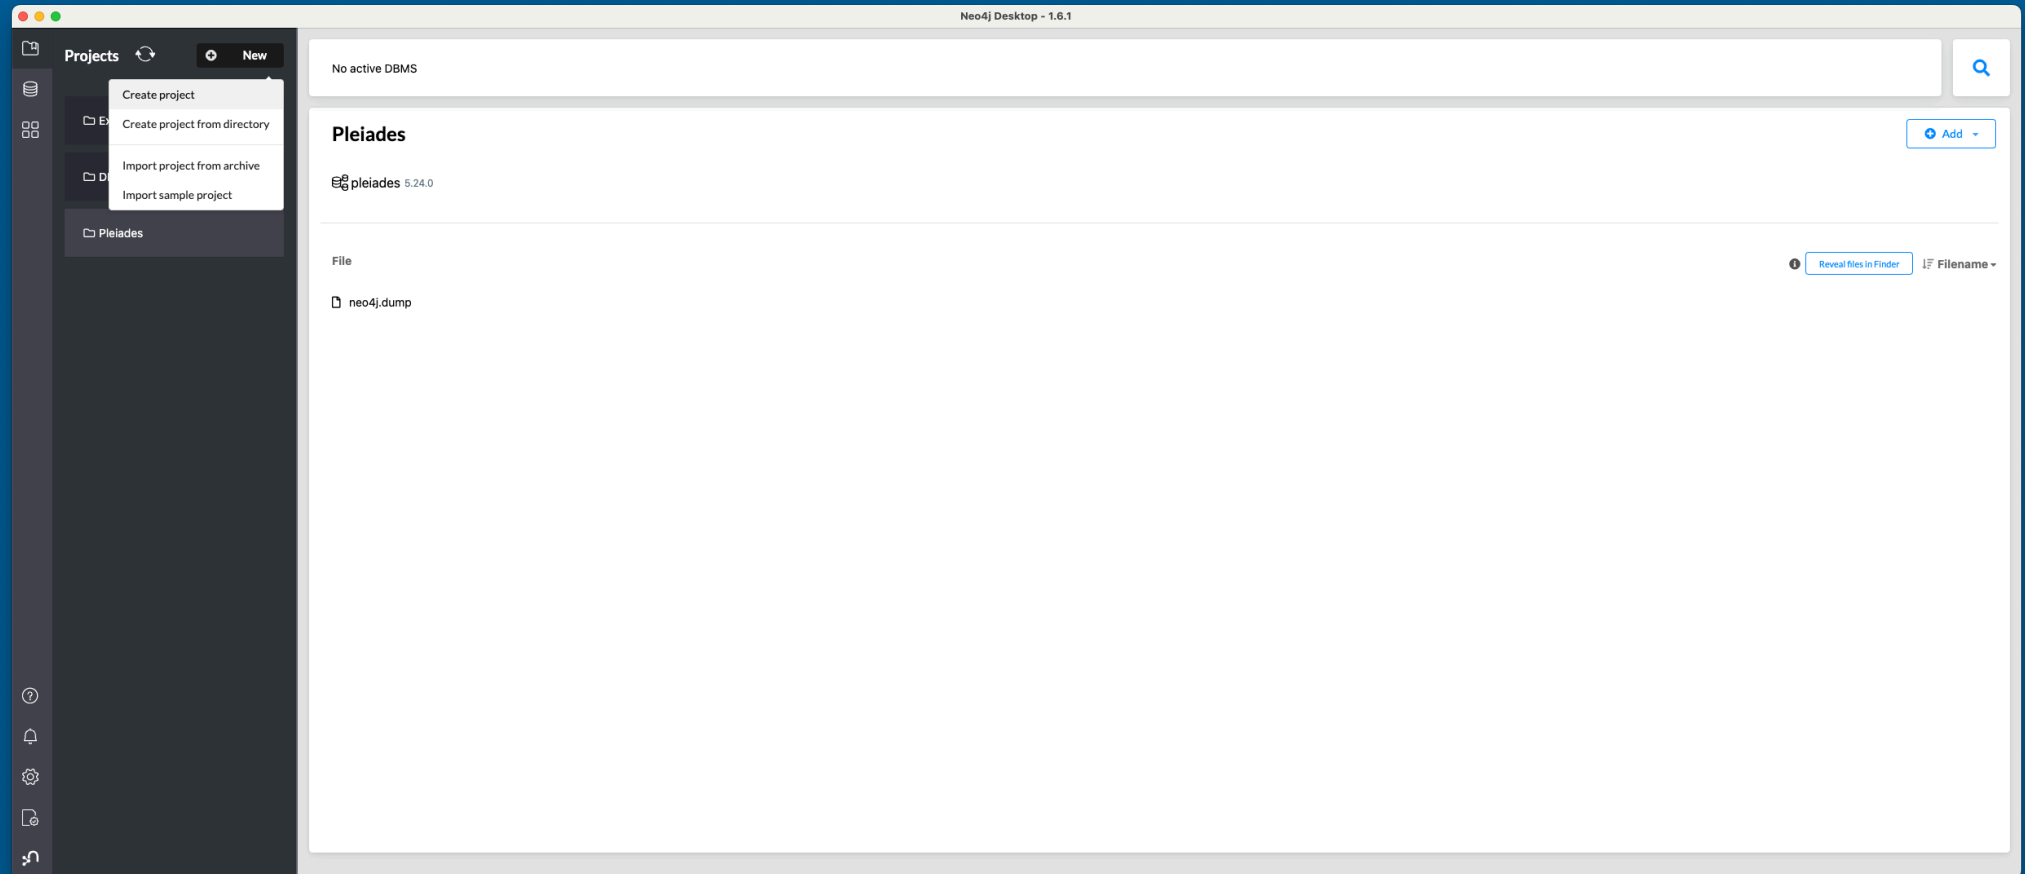
Basic Screen

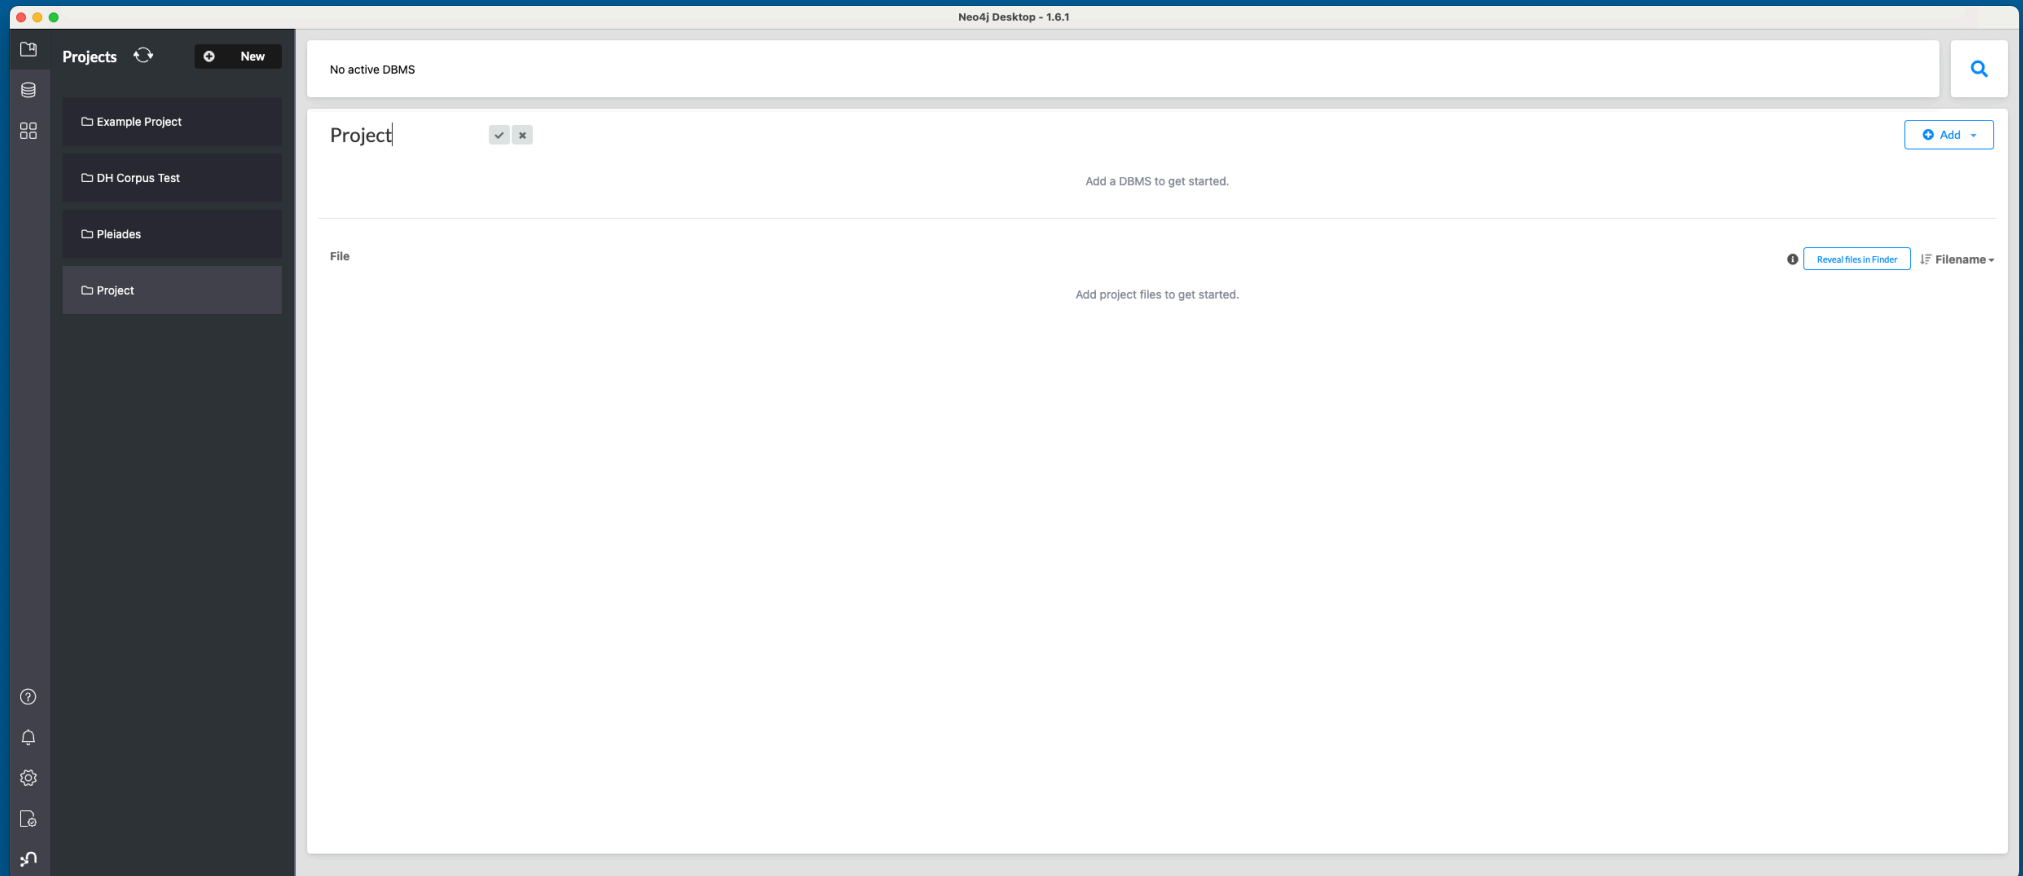


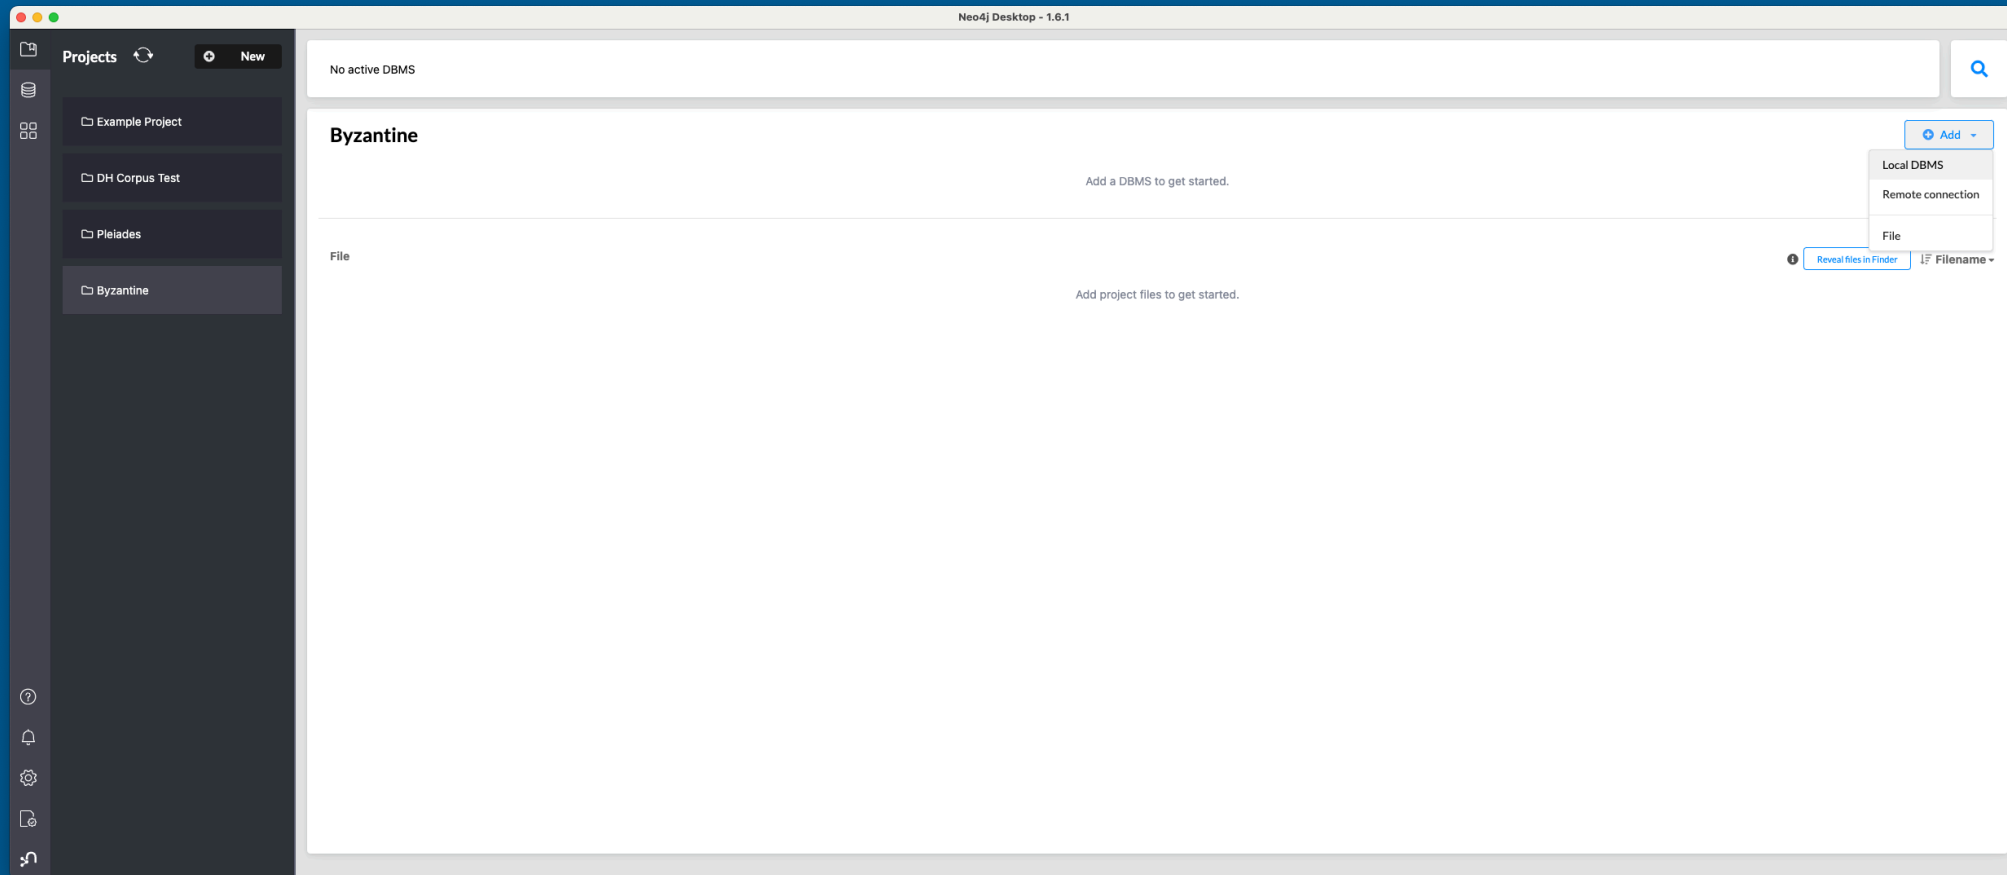
Create a New Project

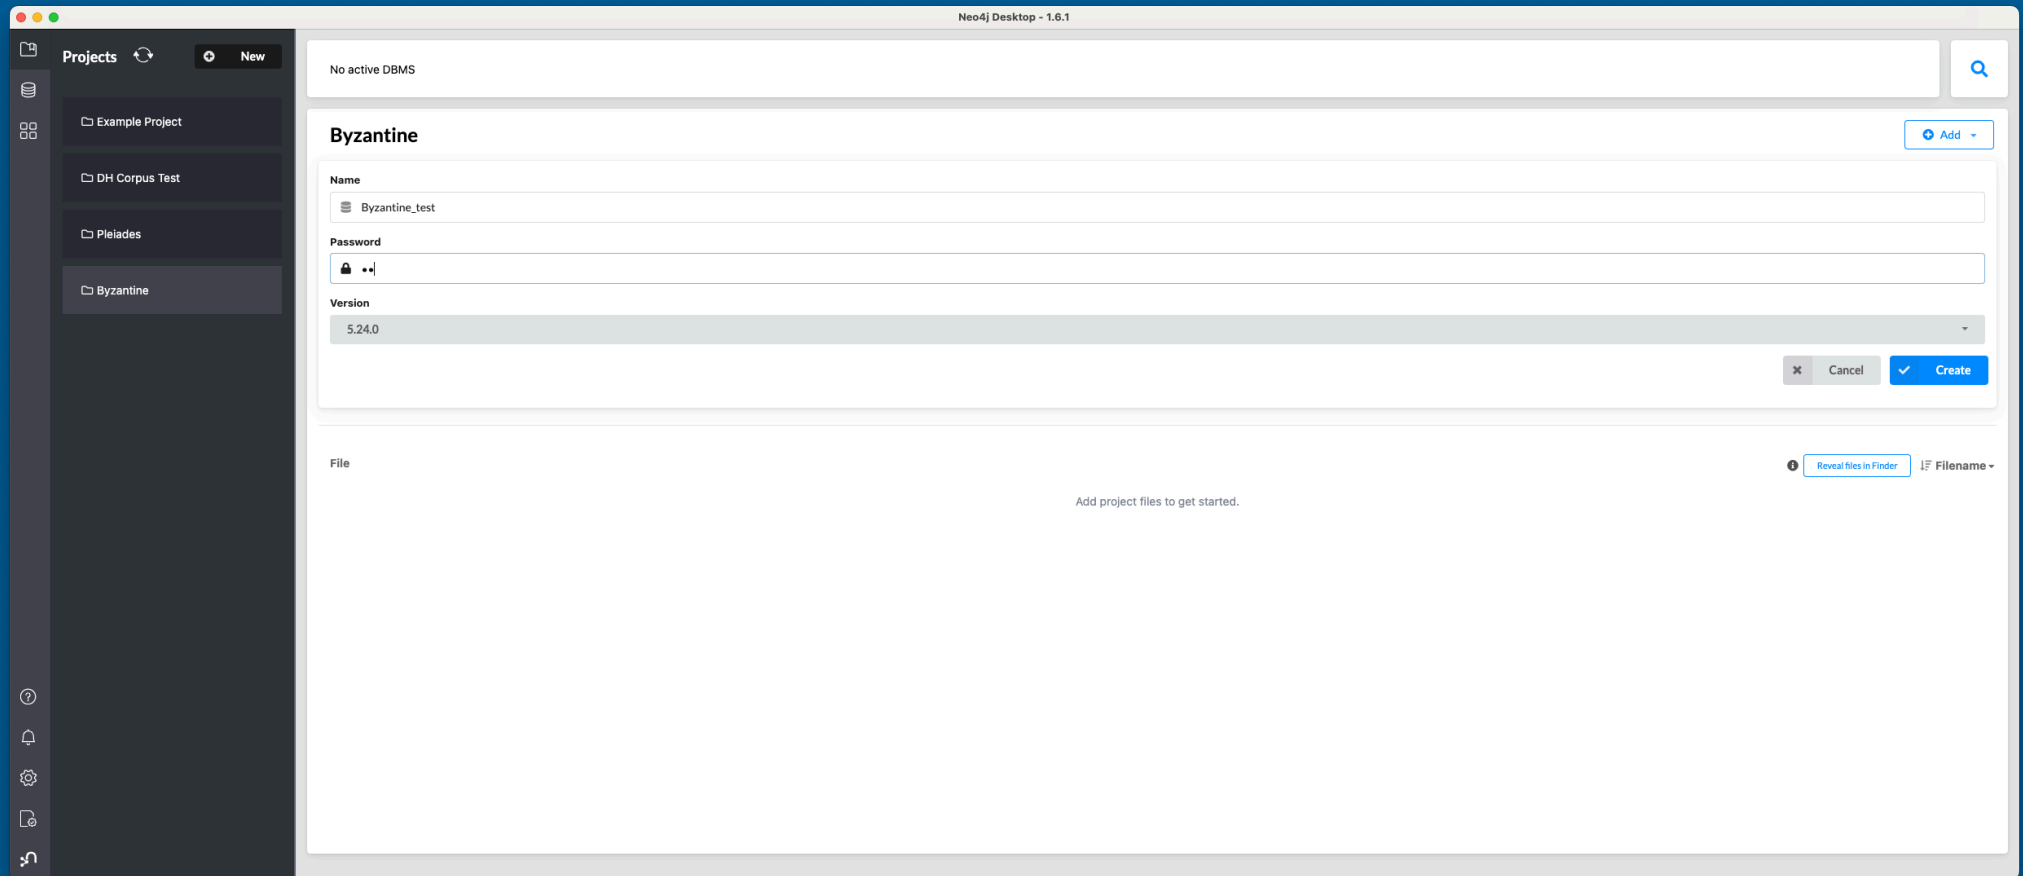


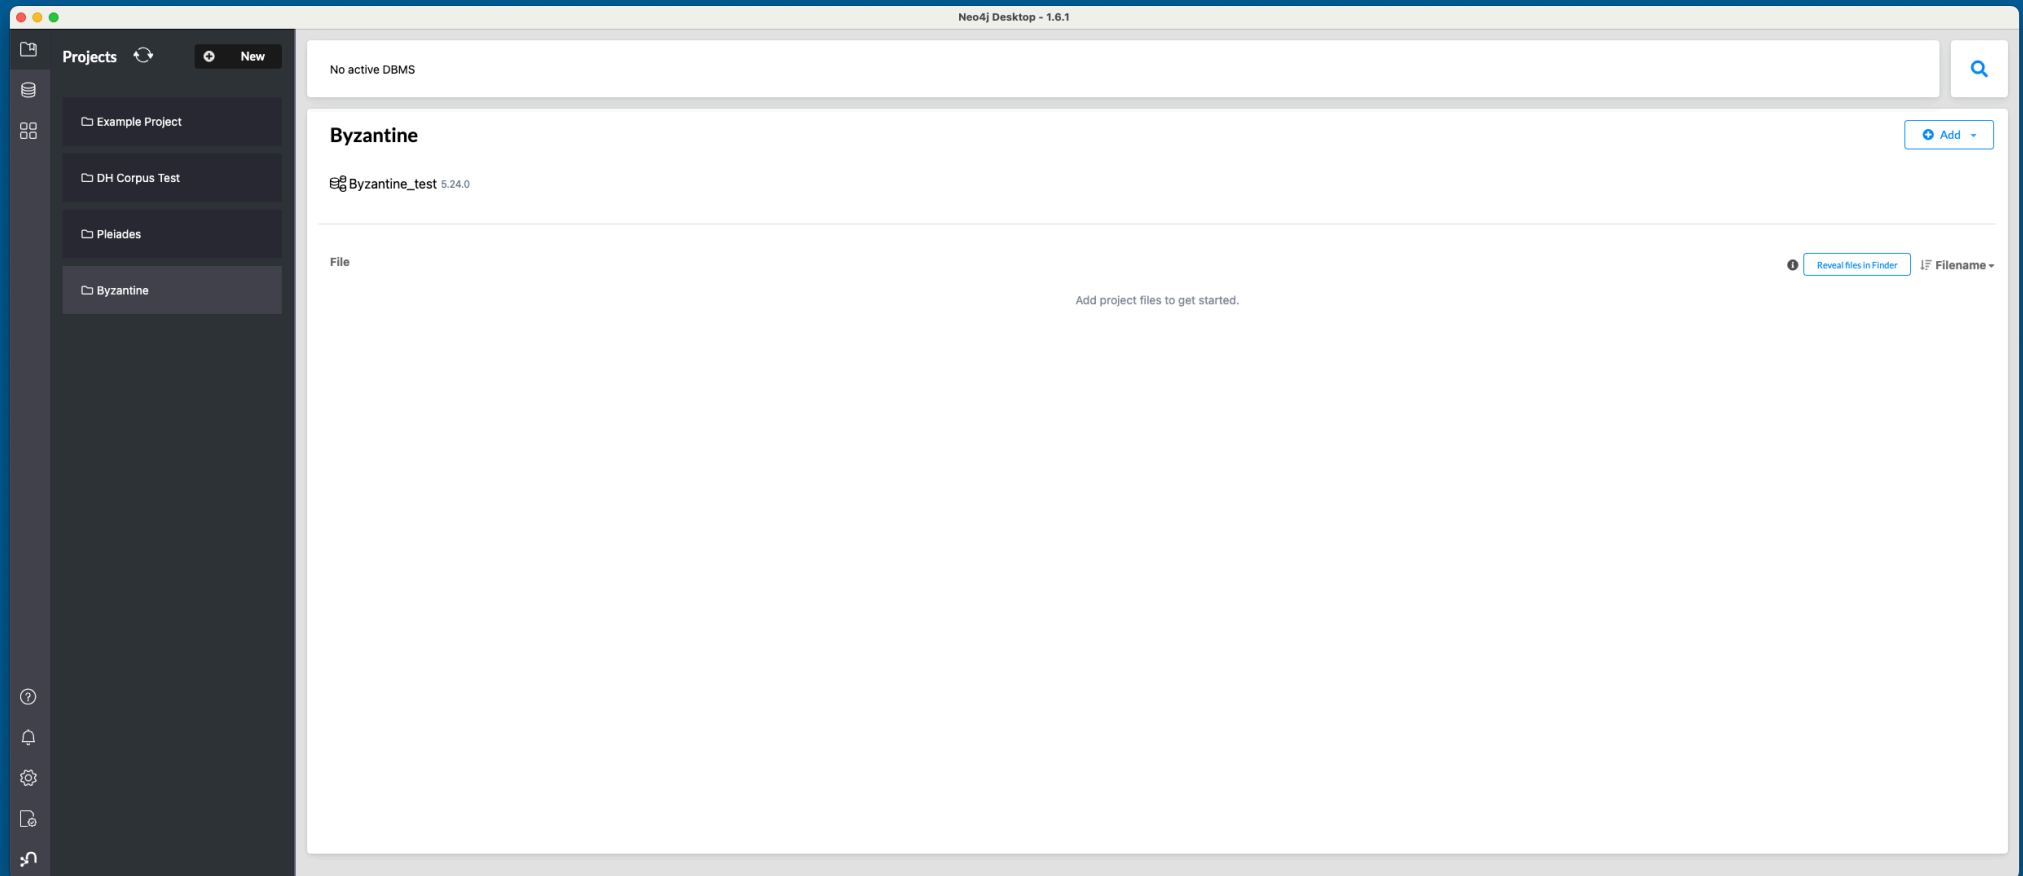
Create a New Project

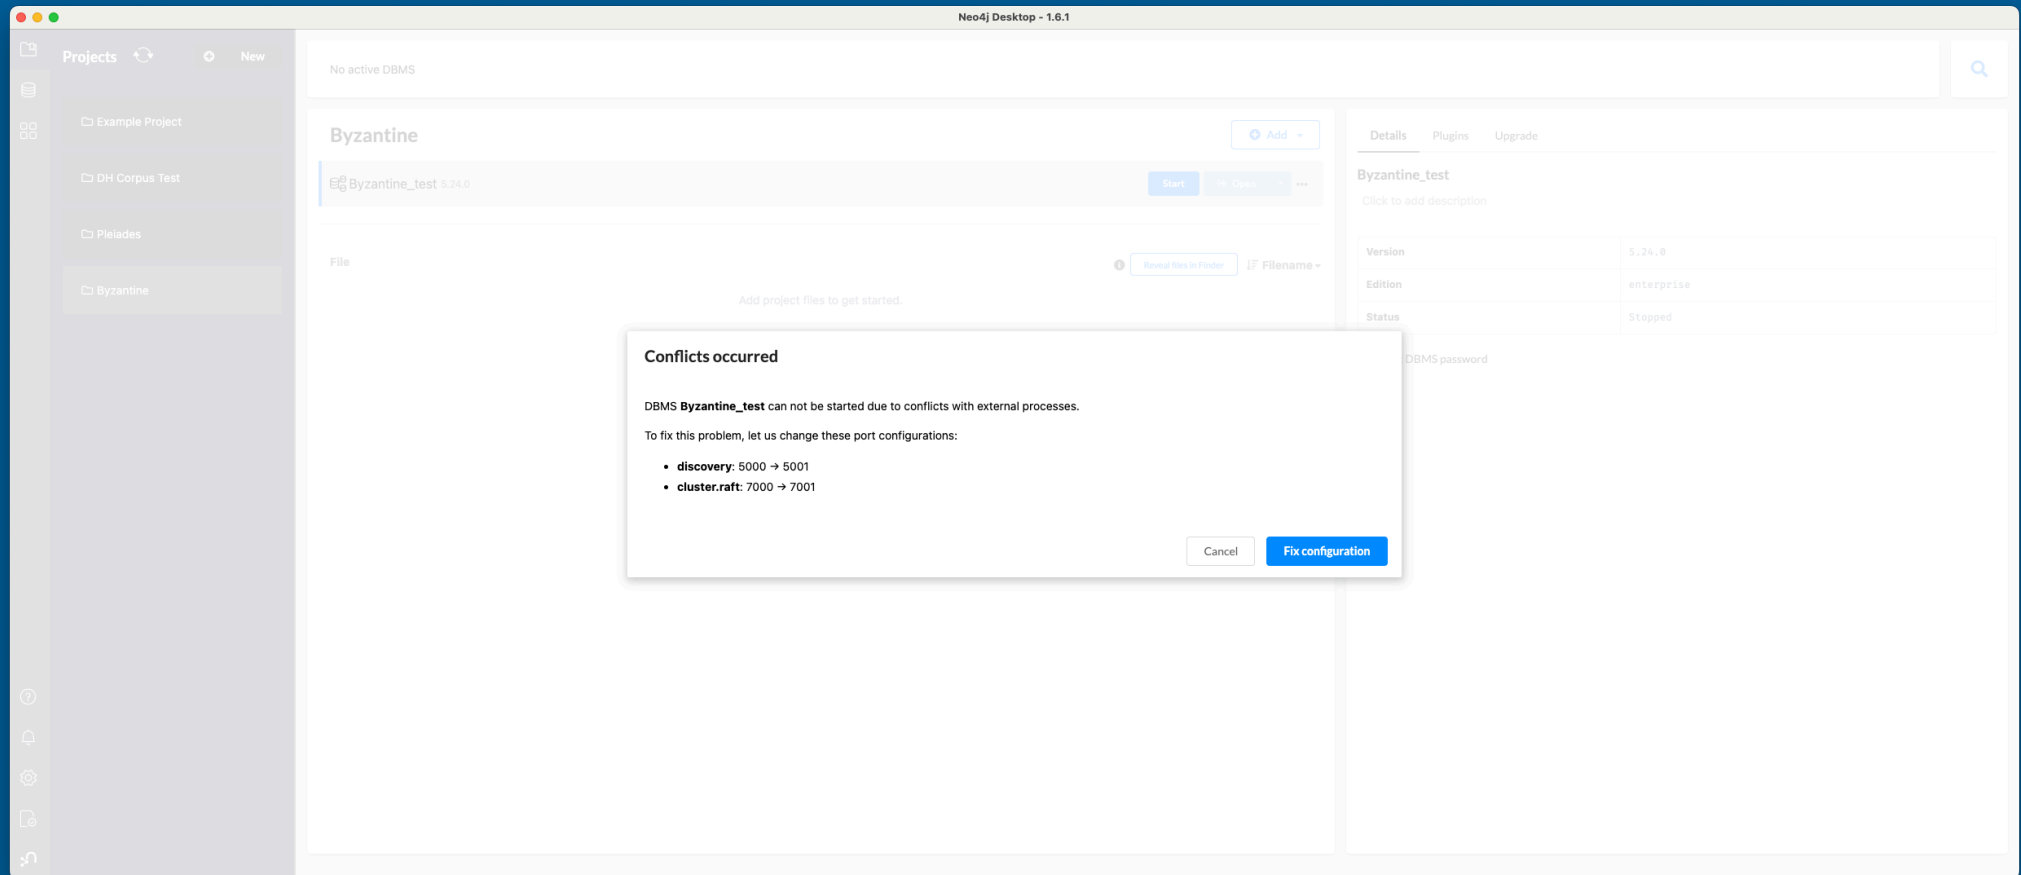












Neo4j Desktop - 1.6.1

Projects

New

Example Project

DH Corpus Test

Pleiades

Byzantine

Byzantine

Byzantine_test 5.24.0

system

neo4j (default)

Create database Refresh

File

Add project files to get started.

Reveal files in Finder

Filename

Stop Open

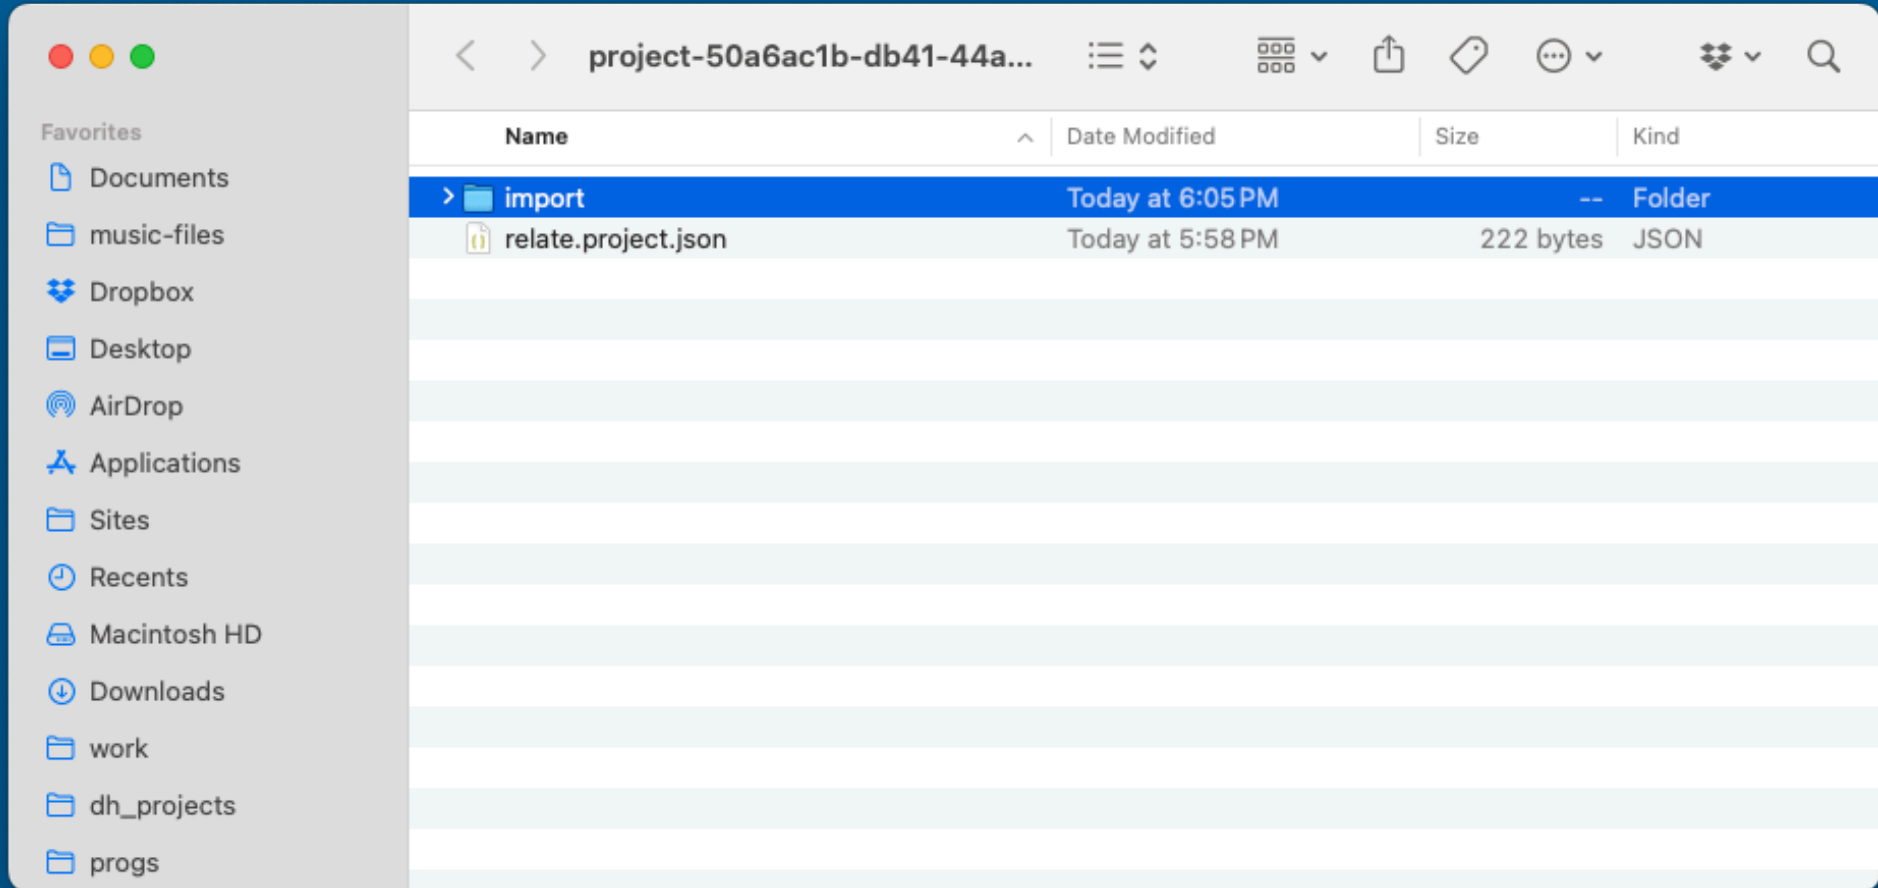
Details Plugins Upgrade

Byzantine_test

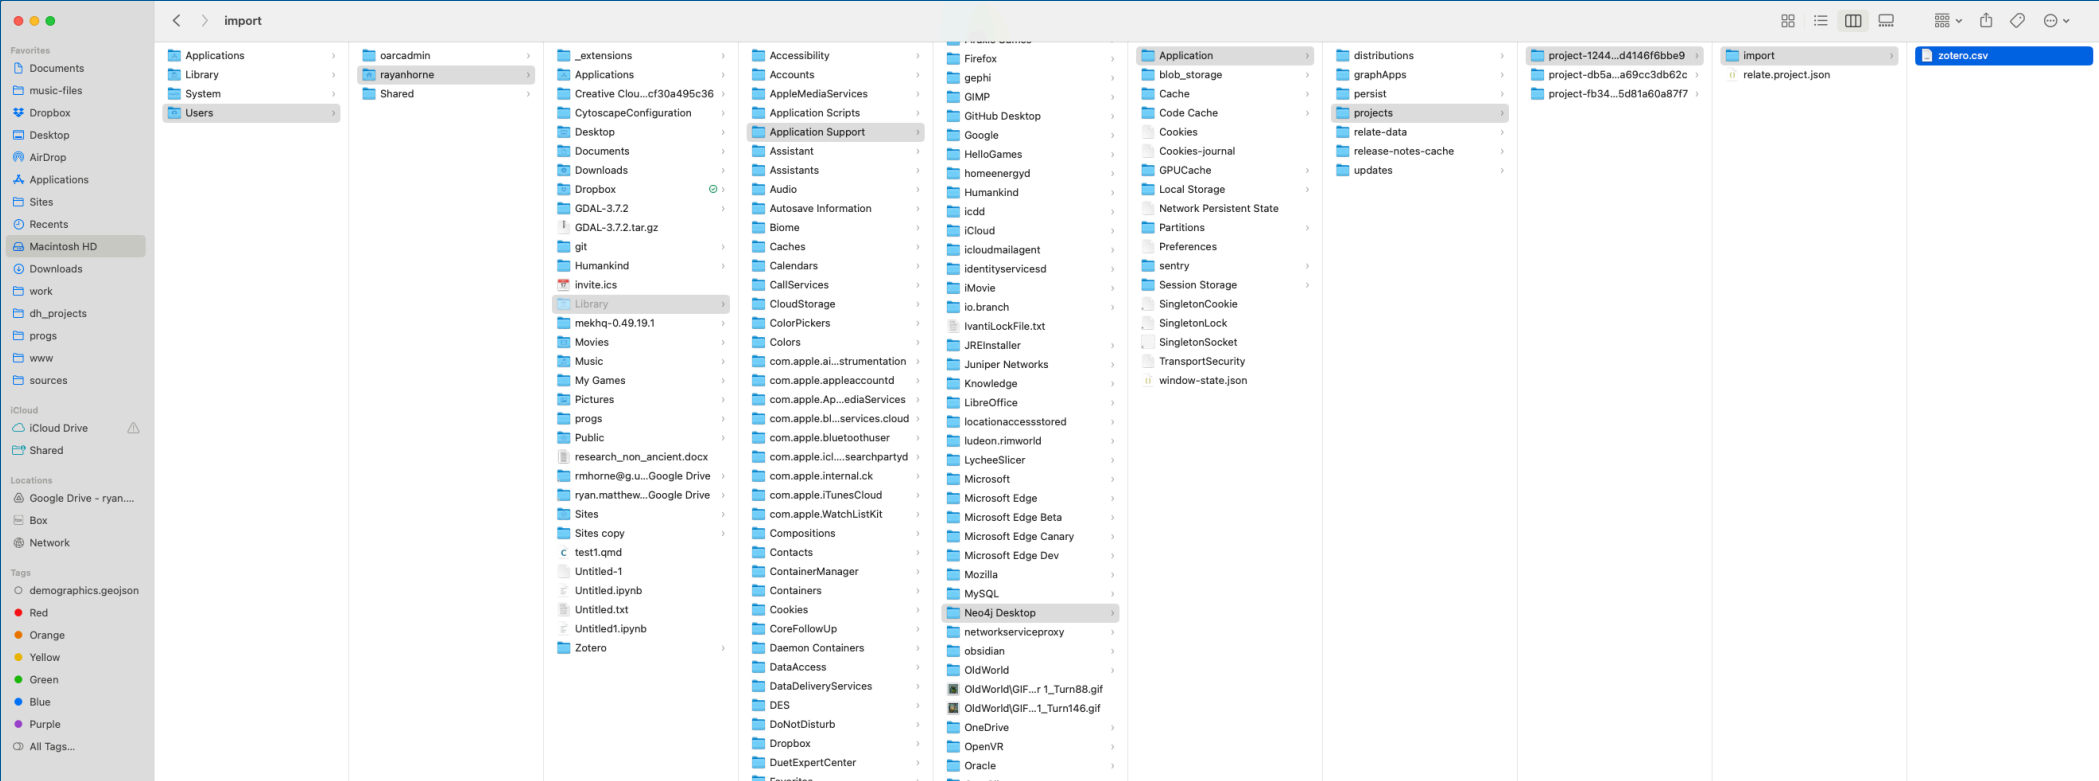
Click to add description

Version	5.24.0
Edition	enterprise
Status	Active
IP address	localhost
Bolt port	7687
HTTP port	7474
HTTPS port	7473

Reset DBMS password



Step 2: Path



Importing CSV Files into Neo4j

Step 1: Prepare Your Data

Structure your CSV with headers like name, id, birth, relation_id

Example: people.csv

```
1 id,name,birth
2 1,Anna Komnene,1083
3 2,John II Komnenos,1087
```

Step 2: Place CSV in import Folder

For Neo4j Desktop: use the built-in import directory

- File path: neo4j/import/people.csv

Step 3: Load Nodes with Cypher

```
1 LOAD CSV WITH HEADERS FROM 'file:///people.csv' AS row
2 CREATE (:Person {id: row.id, name: row.name, birth: toInteger(row.b
```

- **LOAD CSV WITH HEADERS** reads the file
- **CREATE** makes new nodes with properties

Step 4: Load Relationships

Assume a file `wrote.csv`:

```
1 author_id,text_title
2 1,Alexiad
3 2,History
```

```
1
2 LOAD CSV WITH HEADERS FROM 'file:///wrote.csv' AS row
3 MATCH (p:Person {id: row.author_id})
4 CREATE (p)-[:WROTE]->(:Text {title: row.text_title})
```

Step 5: Check Results

```
1  
2 MATCH (p:Person)-[:WROTE]->(t:Text)  
3 RETURN p.name, t.title
```

Thank You!

Any Questions?